

Intrinsic Dimension Estimation as a Design Principle for Autoencoder Bottleneck Dimensions

Anonymous Author

Abstract—Selecting the bottleneck dimension m of an autoencoder (AE) has traditionally relied on expensive manual sweeps or empirical heuristics, leaving the design rationale unclear. This paper proposes a two-stage pipeline that uses intrinsic dimension (ID) estimation as a principled guide for bottleneck design: (1) estimate the intrinsic dimension $\hat{d}_{\text{ID}}$ via TwoNN/MLE to narrow the search range, and (2) identify the optimal bottleneck dimension m^* using the active units (AU) pattern of a variational autoencoder (VAE)—adopting the saturation dimension m^{sat} for manifolds with sharp saturation, or the AU inflection point m^{knee} for gradually increasing datasets—with CKA and MSE elbow as auxiliary criteria. Experiments on synthetic manifolds (Swiss Roll, Torus, Möbius Band) and MNIST yield three main findings: (RQ1) $\hat{d}_{\text{ID}}$ serves as a strong prior indicator, aligning well with m^* especially on synthetic manifolds (with limited accuracy on real data); (RQ2) For topologically nontrivial manifolds (Torus, Möbius Band), the VAE AU saturation value corresponds to the topological embedding dimension d_{emb} under standardized preprocessing; (RQ3) Decoder Jacobian regularization (IsometricAE) achieves near-isometric embedding ($\kappa \approx 1$), realizing the prerequisite for reliable ID estimation in latent space. A secondary finding quantifies the temporal separation of global and local structure formation during MNIST training.

Index Terms—Autoencoder, bottleneck dimension, intrinsic dimension, variational autoencoder, active units, isometric embedding, manifold hypothesis, TwoNN, centered kernel alignment.

I. INTRODUCTION

THE manifold hypothesis [1], [2] posits that high-dimensional natural data concentrate near a smooth low-dimensional manifold $\mathcal{M}$ (with $\dim \mathcal{M} = d \ll N$) embedded in $\mathbb{R}^N$. Autoencoders (AEs) are implicitly designed to capture this manifold in latent space, yet the bottleneck dimension m is typically chosen by manual validation sweeps—a computationally expensive process with no principled justification.

A. Research Aim

We propose to *estimate the intrinsic dimension d_{ID} of the data in advance* and use it as a design guide for the bottleneck dimension. We adopt TwoNN [3] and MLE [4] as ID estimators and quantify their relationship to the optimal AE bottleneck dimension.

Theoretical motivation. Two distinct mechanisms link d_{ID} and m^* :

- **Why AEs need at least d_{ID} dimensions:** When data lie locally on a d_{ID} -dimensional manifold, the bottleneck must satisfy $m \geq d_{\text{ID}}$ to bound reconstruction error under standard smoothness assumptions (Theorem 1).

Hence d_{ID} provides a principled *lower-bound guideline* estimable before training.

- **Why VAE AU converges near d_{emb} :** The support $\mathbb{R}^m$ of the Gaussian prior in a VAE is topologically contractible; embedding a manifold with nontrivial fundamental group π_1 (e.g., Torus with $\pi_1 = \mathbb{Z}^2$) without self-intersection requires $d_{\text{emb}} > d_{\text{ID}}$ dimensions [5]. Under KL regularization, AU tends to converge near d_{emb} in sufficiently capable VAEs (Proposition 3).

B. Two Notions of Dimension

This paper employs two fundamentally different concepts of manifold dimension:

Local Intrinsic Dimension d_{ID} : The dimension of the tangent space at each point of $\mathcal{M}$ —a local geometric quantity measuring the degrees of freedom in a neighborhood. Both the Swiss Roll and the Torus have $d_{\text{ID}} = 2$. TwoNN and MLE estimate d_{ID} from local distance relations.

Topological Embedding Dimension d_{emb} : The minimum dimension of Euclidean space $\mathbb{R}^{d_{\text{emb}}}$ into which $\mathcal{M}$ can be continuously embedded without self-intersection. For simply connected manifolds ($\pi_1 = \{1\}$, e.g., Swiss Roll) $d_{\text{emb}} = d_{\text{ID}}$; for manifolds with nontrivial π_1 (e.g., Torus: $\pi_1 = \mathbb{Z}^2$) $d_{\text{emb}} > d_{\text{ID}}$ [5], [6].

Central claim: AEs reflect d_{ID} (tangent-space dimension) through MSE minimization, while VAE AU saturation reflects d_{emb} (topological embedding dimension) via the continuous-mapping constraint. These two findings coexist without logical contradiction.

C. Research Questions

- RQ1** Can TwoNN/MLE accurately estimate d_{true} on synthetic manifolds, and does $\hat{d}_{\text{ID}}$ align with m^* to a degree useful for design guidance?
- RQ2** To what extent does VAE AU_{sat} reflect the topological properties (fundamental group π_1) of the manifold?
- RQ3** Can isometric embedding be achieved by decoder Jacobian regularization, and how does this differ from encoder-side regularization (CAE)?

D. Contributions

- 1) **Systematic evaluation of ID estimators (RQ1):** We benchmark TwoNN and MLE on synthetic manifolds and demonstrate that $\hat{d}_{\text{ID}}$ serves as a strong prior for AE bottleneck design, particularly for synthetic data (Experiments E1–E3).

- 2) **Observation of VAE AU_{sat} on topologically non-trivial manifolds (RQ2):** On Torus and Möbius Band, VAE AU_{sat} corresponds to the topological minimum embedding dimension under standardized preprocessing (Experiments EA, EB).
- 3) **Isometric embedding via IsometricAE (RQ3):** Decoder Jacobian regularization achieves $\kappa \approx 1$ (Experiment EC).
- 4) **Temporal separation of global/local structure in MNIST (secondary finding):** We quantify how global neighborhood structure forms rapidly in early training while local ID increases gradually (Experiment E6).

II. RELATED WORK

A. Intrinsic Dimension Estimation

TwoNN [3] estimates ID from the ratio $\mu_i = r_2/r_1$ of two nearest-neighbor distances, offering high accuracy and computational efficiency. MLE [4] constructs a maximum-likelihood estimator from k -nearest-neighbor distances, giving stable estimates across choices of k . DANCo [11] uses joint distributions of angles and distances for finer estimation at higher computational cost. Pope et al. [12] showed that natural images have much lower ID than pixel dimension as measured by TwoNN. More recent work has analyzed ID of deep network representations: Ansuini et al. [13] quantified layer-wise ID in deep networks; Birdal et al. [14] studied the connection between ID, persistent homology, and generalization; Valeriani et al. [15] reported a “hump” structure of layer-wise ID in large transformer models.

A known limitation of TwoNN is its *underestimation bias* near manifold boundaries and in high-curvature regions [3]. ID estimates at $m = d_{\text{true}}$ (the boundary of the underfitting regime) are particularly susceptible to this bias.

B. VAE Gaussian Prior and Topological Embedding

The support of the standard Gaussian prior $\mathcal{N}(0, I_m)$ is topologically contractible (any loop can be continuously contracted to a point) [5]. Therefore, embedding a manifold with nontrivial π_1 (e.g., Torus T^2 with $\pi_1 = \mathbb{Z}^2$) without self-intersection requires at least $d_{\text{emb}} > d_{\text{ID}}$ dimensions [5], [6]. KL regularization thus tends to drive AU_{sat} toward d_{emb} in sufficiently capable VAEs (Proposition 3, Experiments EA, EB).

C. AE Bottleneck Dimension and Representation Quality

β -VAE [16] promotes disentangled representations by increasing the weight of KL regularization. Lucas et al. [17] analyzed posterior collapse conditions in linear VAEs and theoretically showed that AU saturates near d_{ID} ; however, their analysis assumes a single linear manifold and does not systematically examine topologically nontrivial cases ($AU_{\text{sat}} \approx d_{\text{emb}}$). Nazari et al. [10] and Kato et al. [9] proposed isometric AEs via decoder Jacobian regularization. Contractive AE (CAE) [18] penalizes the encoder Frobenius norm; denoising AE (DAE) [19] implicitly estimates the score function and learns tangent-direction-selective representations.

Distinction from prior work: Existing studies mostly evaluate AE performance with a single metric. This paper constructs a multi-metric pipeline (ID estimation, AU, CKA, condition number κ , Trustworthiness) and demonstrates $m^* \approx d_{\text{ID}}$ on both synthetic and real data, with explicit consideration of TwoNN’s underestimation bias.

III. PROPOSED FRAMEWORK

We denote the bottleneck dimension consistently as m ($:= d_z$), the optimal bottleneck as m^* , and the AU saturation dimension as m^{sat} .

A. Notation Summary

Table I summarizes the key symbols used throughout this paper.

B. Two-Stage Pipeline

The proposed framework consists of two stages.

Stage 1: ID Estimator Evaluation and Selection. Benchmark TwoNN and MLE on synthetic manifolds with known d_{true} to select a reliable estimator. For real data with large input dimension N , apply TwoNN/MLE to the latent space of a reference AE trained with a large bottleneck ($m_{\text{ref}} \gg \hat{d}_{\text{ID}}$).

Stage 2: Optimal Bottleneck Search via VAE Sweep. Sweep m over a discrete grid $\mathcal{M} = \{1, 2, 4, 8, 12, 16, 20, 32, 64, 128, 256\}$, training a VAE ($\beta = 4$) at each m and computing $AU(m)$.

Remark 1 (Grid Resolution). *Since $\mathcal{M}$ is an approximately log-scale grid, when d_{ID} falls between grid points (e.g., $d_{\text{ID}} = 6$ between $m = 4$ and $m = 8$), the grid point where AU first exceeds d_{ID} serves as a conservative upper bound for m^* . Practical recommendation: verify with a finer grid near $\hat{d}_{\text{ID}}$.*

Fig. 1 illustrates the proposed two-stage pipeline.

C. Optimal Bottleneck Dimension

The normalized AU growth rate is defined as

$$r(m) = \frac{AU(m) - AU(m_{\text{prev}})}{m - m_{\text{prev}}} \quad (1)$$

where m_{prev} is the preceding grid point in $\mathcal{M}$.

AU Pattern Classification:

- **Saturation-type:** AU maintains the same value at two or more consecutive grid points (plateau width ≥ 2). Typical examples: Swiss Roll, Torus.
- **Gradual-type:** AU increases nearly monotonically across the entire grid; $AU(m_{\text{max}}) \approx AU_{\text{sat}}$ is reached only at the largest m . Typical example: MNIST.

Case-based definition of m^* :

$$m^* = \begin{cases} m^{\text{sat}} := \min\{m \in \mathcal{M} \mid AU(m) = AU_{\text{sat}}\} & (\text{saturation-type}) \\ m^{\text{knee}} := \min\{m \in \mathcal{M} \mid r(m) < \tau_{\text{knee}} \cdot r_{\text{max}}\} & (\text{gradual-type}) \end{cases} \quad (2)$$

TABLE I
KEY NOTATION

Symbol	Meaning	Reference
m	AE bottleneck dimension (design parameter)	§III
m^*	Optimal bottleneck dimension	§III
m^{sat}	Dimension at which AU first saturates (saturation-type)	Eq. (3)
m^{knee}	AU growth inflection point (gradual-type)	Eq. (4)
AU_{sat}	Maximum active units reached on grid	Eq. (3)
$\text{AU}(m)$	Active unit count at bottleneck m	§III
$\hat{d}_{\text{ID}}$	ID estimate from TwoNN/MLE	§III
d_{true}	True ID of synthetic manifold (known)	§V-B
d_{emb}	Topological minimum embedding dimension	§I-B
κ	Condition number of decoder Jacobian (isometry metric)	§II
$r(m)$	Normalized AU growth rate	Eq. (1)
τ_{knee}	Inflection-point threshold (0.25 in this work)	Eq. (4)




Fig. 1. Proposed two-stage pipeline for bottleneck dimension design. Stage 1 trains a reference AE to estimate $\hat{d}_{\text{ID}}$ via TwoNN/MLE. Stage 2 sweeps a VAE over a grid, classifies the AU pattern as saturation-type or gradual-type, and selects m^* using CKA and MSE as auxiliary criteria.

where

$$\text{AU}_{\text{sat}} = \max_{m' \in \mathcal{M}} \text{AU}(m'), \quad (3)$$

$$r_{\text{max}} = \max_{m' \in \mathcal{M}} r(m'), \quad \tau_{\text{knee}} = 0.25. \quad (4)$$

Design Rule (two-branch):

Saturation-type: Given the saturation set $\mathcal{C}_{\text{sat}} = \{m \in \mathcal{M} \mid \text{AU}(m) = \text{AU}_{\text{sat}}\}$,

$$m^* = \min \left\{ m \in \mathcal{C}_{\text{sat}} \mid \text{CKA}(m) \geq (1 - \varepsilon) \max_{m' \in \mathcal{C}_{\text{sat}}} \text{CKA}(m') \right\} \quad (5)$$

with $\varepsilon = 0.01$. Verify consistency with the MSE elbow m_{elbow}^* (Eq. (12)).

Gradual-type: Given $\mathcal{C}_{\text{knee}} = \{m \in \mathcal{M} \mid m \geq m^{\text{knee}}, r(m) < \tau_{\text{knee}} \cdot r_{\text{max}}\}$,

$$m^* = \min \left\{ m \in \mathcal{C}_{\text{knee}} \mid \text{CKA}(m) \geq (1 - \varepsilon) \max_{m' \in \mathcal{C}_{\text{knee}}} \text{CKA}(m') \right\} \quad (6)$$

Rationale for AU as primary criterion: Unlike AIC/BIC, which require an effective parameter count (ill-defined for nonlinear AEs), AU saturation directly measures the *information dimensions actually used* by the model under KL regularization, providing a data-driven model selection criterion independent of AIC/BIC (see Proposition 3).

D. Evaluation Metrics

TwoNN ID estimator [3]: Uses the ratio $\mu_i = r_2/r_1$ of the two nearest-neighbor distances,

$$\hat{d}_{\text{ID}}^{\text{TwoNN}} = \frac{1}{\mathbb{E}[\log \mu_i]}. \quad (7)$$

MLE ID estimator [4]: Maximum-likelihood estimate from k -nearest-neighbor distance ratios,

$$\hat{d}_{\text{ID}}^{\text{MLE}} = \left[\frac{1}{k-1} \sum_{j=1}^{k-1} \log \frac{r_k(\mathbf{x}_i)}{r_j(\mathbf{x}_i)} \right]^{-1}. \quad (8)$$

Active Units (AU) [17]: A latent dimension k is active if $V_k = \text{Var}_{\mathbf{x}}[\mathbb{E}_{q_\phi}[z_k|\mathbf{x}]] > \delta$,

$$\text{AU} = \sum_{k=1}^m \mathbf{1}[V_k > \delta], \quad \delta = 10^{-2}. \quad (9)$$

CKA [20]: Given input $\mathbf{X}$ and reconstruction $\hat{\mathbf{X}}$, with kernel matrices $K = \mathbf{X}\mathbf{X}^\top$, $L = \hat{\mathbf{X}}\hat{\mathbf{X}}^\top$, and centering matrix $H = I_n - \frac{1}{n}\mathbf{1}\mathbf{1}^\top$,

$$\text{CKA}(K, L) = \frac{\text{HSIC}(K, L)}{\sqrt{\text{HSIC}(K, K) \cdot \text{HSIC}(L, L)}} \quad (10)$$

where $\text{HSIC}(K, L) = \frac{1}{(n-1)^2} \text{tr}(KHLH)$.

Condition number κ : Ratio of largest to smallest singular values of decoder Jacobian $J_g(z) \in \mathbb{R}^{N \times m}$, $\kappa = \sigma_1/\sigma_m$.

MSE elbow: After smoothing $\tilde{\mathcal{L}}(m_k) = \frac{1}{2w+1} \sum_{j=-w}^w \mathcal{L}(m_{k+j})$ (window $w = 1$), we compute normalized second differences

$$\Delta^2(m_k) = \frac{\tilde{\mathcal{L}}(m_{k-1}) - 2\tilde{\mathcal{L}}(m_k) + \tilde{\mathcal{L}}(m_{k+1}))}{\tilde{\mathcal{L}}(m_1)}, \quad (11)$$

$$m_{\text{elbow}}^* = \arg \max_{k \in \{2, \dots, K-1\}} \Delta^2(m_k). \quad (12)$$

IV. THEORETICAL BACKGROUND

The following theorems and propositions provide the theoretical interpretation of our experiments. Proofs are included for all results.

Theorem 1 (Bottleneck Lower Bound). *For a data distribution $p(\mathbf{x})$ concentrated on a smooth d -dimensional manifold $\mathcal{M} \subset \mathbb{R}^N$, having bounded expected reconstruction error requires $m \geq d$.*

Proof. Consider an AE with encoder $f : \mathbb{R}^N \rightarrow \mathbb{R}^m$ and decoder $g : \mathbb{R}^m \rightarrow \mathbb{R}^N$. Assume $m < d$. Since f is continuous, the image of the composition $g \circ f : \mathbb{R}^N \rightarrow \mathbb{R}^N$ is contained in $g(\mathbb{R}^m)$, a set of dimension at most m . Since $\mathcal{M}$ is a smooth d -dimensional manifold with $d > m$, by Sard's theorem the intersection of $\mathcal{M}$ with any set of dimension $\leq m$ has $\mathcal{M}$ -measure zero. Therefore $g(f(\mathbf{x})) = \mathbf{x}$ cannot hold p -almost everywhere, so $\mathbb{E}_p[\|\mathbf{x} - g(f(\mathbf{x}))\|^2] > 0$ cannot be made to vanish. Hence $m \geq d$ is necessary. $\square$

Experiment E1 (Swiss Roll, $d_{\text{true}} = 2$) confirms this: MSE drops approximately 35-fold ($1.016 \rightarrow 0.029$, 3-seed average) at the $m = 1 \rightarrow 2$ transition.

Proposition 1 (Effective Rank of Decoder Jacobian). *When a sufficiently capable AE achieves accurate reconstruction of $\mathcal{M}$, the number of significant singular values of $J_g(z) \in \mathbb{R}^{N \times m}$ equals $r^* = \min(m, d_{\text{emb}})$. For $m \leq d_{\text{emb}}$: $\text{rank } J_g(z) = m$; for $m > d_{\text{emb}}$: surplus dimensions degenerate and $\text{rank } J_g(z) = d_{\text{emb}}$.*

Proof. When the AE accurately encodes $\mathcal{M}$, the image of g locally covers a neighborhood of $\mathcal{M}$. The image of the differential $J_g(z) : \mathbb{R}^m \rightarrow \mathbb{R}^N$ spans the tangent space $T_{g(z)}\mathcal{M}$. Since $\mathcal{M}$ is d -dimensional, $\dim T_{g(z)}\mathcal{M} = d$, so $\text{rank}(J_g(z)) = \min(m, d)$. When $m > d$, the $m - d$ surplus directions in $\mathbb{R}^m$ have no corresponding tangent-space components; gradient-based training drives their singular values toward zero. $\square$

Experiment E2 confirms: for $m > d_{\text{emb}}$ (Torus, $m = 4$, $d_{\text{emb}} = 3$), the significant singular value count is 3.00.

Theorem 2 (Pullback Metric and Isometry). *When decoder $g : \mathbb{R}^m \rightarrow \mathbb{R}^N$ is an isometric embedding into the Riemannian manifold $(\mathcal{M}, g_{\mathcal{M}})$, the pullback metric $G(z) = J_g(z)^\top J_g(z)$ is proportional to the identity, i.e., $\kappa = 1$.*

Proof. By definition of an isometric embedding, for any $z \in \mathbb{R}^m$ and tangent vectors $u, v \in \mathbb{R}^m$: $(g_{\mathcal{M}})_{g(z)}(J_g(z)u, J_g(z)v) = u^\top v$. The pullback metric satisfies $(g^*g_{\mathcal{M}})_z(u, v) = u^\top G(z)v$, where $G(z) = J_g(z)^\top J_g(z)$. Isometry therefore implies $u^\top G(z)v = u^\top v$ for all u, v , so

$G(z) = I_m$. Hence all singular values of $J_g(z)$ are equal to 1, giving $\kappa = \sigma_{\text{max}}/\sigma_{\text{min}} = 1$. $\square$

Experiment EC shows IsometricAE achieves $\kappa \approx 1.055$.

Proposition 2 (Noise Selectivity). *The optimal denoising AE approximates the score function $\nabla_{\mathbf{x}} \log p(\mathbf{x})$. An AE trained with $m = d_{\text{true}}$ selectively removes noise in the normal direction more than in the tangent direction.*

Proof. By the result of Alain and Bengio [19], the reconstruction residual of a denoising AE satisfies $r(\mathbf{x}) - \mathbf{x} \approx \sigma^2 \nabla_{\mathbf{x}} \log p(\mathbf{x})$. Under the manifold hypothesis, data density $p(\mathbf{x})$ decays sharply in normal directions and varies slowly along tangent directions. Hence $\nabla_{\mathbf{x}} \log p(\mathbf{x})$ points predominantly in the normal direction, and the reconstruction correction selectively suppresses normal-direction noise. $\square$

Experiment E3 confirms: at $m = d_{\text{true}}$, normal-direction error $>$ tangent-direction error.

Proposition 3 (KL Regularization Required for AU Saturation). *For a Vanilla AE (no regularization), $\text{AU} = m$ for any $m > d_{\text{ID}}$. AU saturation at d_{ID} requires KL regularization (the information bottleneck constraint).*

Proof. The VAE objective is $\mathcal{L} = \mathcal{L}_{\text{rec}} + \beta D_{\text{KL}}(q(z|\mathbf{x})\|p(z))$. For dimensions beyond the intrinsic dimension of the data, the encoder posterior $q(z_k|\mathbf{x})$ carries no information about $\mathbf{x}$; the KL term forces $q(z_k|\mathbf{x}) \rightarrow p(z_k) = \mathcal{N}(0, 1)$, eliminating the information channel. Consequently, the variance $\text{Var}_{\mathbf{x}}[\mathbb{E}_q[z_k|\mathbf{x}]]$ collapses below the threshold $\delta = 10^{-2}$, deactivating those units. A Vanilla AE has no such penalty, so all m dimensions remain active ($\text{AU} = m$) regardless of the data dimensionality. $\square$

Experiment EA confirms: only VAE ($\beta = 4$) exhibits AU saturation; Vanilla AE always has $\text{AU} = m$.

Proposition 4 (Neighborhood Preservation and Optimal Bottleneck). *At $m = d_{\text{true}}$ (or the topological minimum embedding dimension), Trustworthiness improves sharply, and further increases in m yield limited additional improvement.*

Proof. By Theorem 1, for $m < d$ the encoder f cannot be injective on $\mathcal{M}$, causing unavoidable ‘‘information collisions’’ that distort local neighborhoods, reducing Trustworthiness. At $m \geq d$, the topological obstruction is removed: a homeomorphic (structure-preserving) mapping can be learned, and Trustworthiness improves sharply. For $m > d$, the additional dimensions provide negligible benefit as the neighborhood structure is already well-preserved. $\square$

Experiment E5 confirms: Swiss Roll ($\Delta \approx 0.015$ at $m = 2$) and Torus ($\Delta \approx 0.015$ at $m = 3 = d_{\text{emb}}$).

Table II summarizes the theorem-experiment correspondences.

V. EXPERIMENTS

A. Experimental Setup

Common hyperparameters are listed in Table III; model-specific hyperparameters in Tables IV and V.

TABLE II
THEOREM–EXPERIMENT CORRESPONDENCE

Theorem	Claim	Experiment
Thm. 1 (lower bound)	MSE spikes for $m < d_{\text{true}}$	E1
Prop. 1 (effective rank)	Significant singular values = m	E2
Thm. 2 (pullback metric)	$\kappa \approx 1$ isometry	EC
Prop. 2 (noise selectivity)	Normal direction selectively removed	E3
Prop. 3 (AU saturation)	KL regularization required	EA
Prop. 4 (neighborhood preservation)	Trustworthiness jumps at $m = d_{\text{true}}$ (or d_{emb})	E5
Topological embedding (supplementary)	$\text{AU}_{\text{sat}} \approx d_{\text{emb}}$ (conditional)	EB

TABLE III
COMMON HYPERPARAMETERS

Parameter	Value
Optimizer	Adam ($\text{lr} = 1 \times 10^{-3}$)
Batch size	128
Activation	ReLU (output layer: Sigmoid)
Random seed	42 (fixed)
Framework	PyTorch 2.0

TABLE IV
SYNTHETIC MANIFOLD AE/VAE HYPERPARAMETERS

Parameter	Value
Hidden dims	[64, 32]
Training epochs	200–300
VAE β	4.0
IsometricAE λ_{iso}	0.01
Sample size	5,000 (train 4,000 + test 1,000)

TABLE V
MNIST AE/VAE HYPERPARAMETERS

Parameter	Value
Input dim	784 (28×28 flattened)
Hidden dims	[512, 256]
Training epochs	300 (E4), 200 (E6)
VAE β	4.0
Bottleneck m	1, 2, 4, 8, 12, 16, 20, 32, 64, 128, 256
Batch size	128

ID estimation uses $k = 2$ for TwoNN and $k = 10$ for MLE. The AU threshold $\delta = 10^{-2}$ follows the standard value of [17]. Table VI lists the correspondence between experiments and research questions.

AU threshold sensitivity analysis: We compared AU counts across $\delta \in \{0.001, 0.01, 0.05, 0.1\}$ for Swiss Roll and Torus VAE ($\beta = 4$). As shown in Table VII, AU saturation values are *completely robust* to the choice of δ : Swiss Roll maintains AU= 2 and Torus (at $m \geq 6$) maintains AU= 3 across all δ values. This robustness arises because active dimensions have variance $O(1)$ while inactive dimensions have variance $O(10^{-4})$ or smaller, giving a two-order-of-magnitude gap that is insensitive to threshold changes.

B. ID Estimation and Optimal Bottleneck on Synthetic Manifolds (RQ1)

Question answered (RQ1): Does TwoNN/MLE estimate d_{true} accurately, and does $\hat{d}_{\text{ID}}$ align with m^* as a design




Fig. 2. E1: MSE vs. m (3-seed mean $\pm$ std, log scale). Swiss Roll (solid, $\bullet$): sharp drop at $m = 2 = d_{\text{true}}$. Torus (dashed, $\blacktriangle$): effective elbow at $m = 3$.

guideline?

1) *Reconstruction Error Elbow Analysis (Experiment E1):* We swept m for Swiss Roll ($d_{\text{true}} = 2$, ambient = 3) and Torus ($d_{\text{true}} = 2$, ambient = 3), training AEs and recording MSE, AU, TwoNN ID, and MLE ID (3-seed mean $\pm$ std). Results are shown in Tables VIII and IX.

For Swiss Roll, MSE drops approximately 34-fold ($1.016 \rightarrow 0.029$) at $m = 2 = d_{\text{true}}$, confirming Theorem 1. The standard deviation across 3 seeds is ± 0.001 , demonstrating high robustness. TwoNN ID at $m = 2$ (1.15 ± 0.05) underestimates the true ID of 2; this is the known finite-sample boundary bias [3]. At $m = 3$, TwoNN ID converges to 1.94 ± 0.06 and MLE to 2.09 ± 0.02 , consistent with each other.

For Torus, the effective MSE elbow appears at $m = 3$ despite $d_{\text{true}} = 2$. The topological explanation is discussed in §VI-B.

Fig. 2 shows the MSE elbow curves.

2) *Decoder Jacobian Analysis (Experiment E2):* Decoder Jacobian singular value spectra were evaluated (significant singular value count, condition number κ , and TSA error) as shown in Table X.

For Torus at $m = 4$, the significant singular value count is 3.00 (not 4), confirming Proposition 1: the surplus dimension degenerates to match the topological embedding dimension $d_{\text{emb}} = 3$.

3) *Noise Selectivity (Experiment E3):* We measured the noise selectivity ratio ($\text{NSR} = \text{tangent-direction error} /$

TABLE VI
RESEARCH QUESTIONS AND EXPERIMENT MAPPING

Exp.	RQ	Section	Purpose
E1	RQ1	§V-B	ID estimation accuracy; $m^* \approx d_{\text{ID}}^{\hat{}}$
E2	Prop. 1, Thm. 2 basis	§V-B	Decoder Jacobian rank and κ
E3	Prop. 2	§V-B	Noise selectivity (DAE)
E5	Prop. 4	§V-B4	Trustworthiness vs. m
EA	RQ2	§V-C	VAE AU saturation via KL regularization
EB	RQ2 (depth)	§V-C	AU_{sat} vs. topological embedding dim.
EC	RQ3	§V-C	Isometric embedding achievability
E4	RQ1 (real data)	§V-D	MNIST $m^* \approx d_{\text{ID}}^{\hat{}}$
E6	Novel finding	§V-D	MNIST training dynamics
EX	Statistical reliability	§V-D	Multi-seed stability

TABLE VII
AU THRESHOLD SENSITIVITY (VAE $\beta = 4$): AU VALUES ARE COMPLETELY ROBUST TO THRESHOLD CHOICE

m	Swiss Roll				Torus			
	$\delta=0.001$	$\delta=0.01$	$\delta=0.05$	$\delta=0.1$	$\delta=0.001$	$\delta=0.01$	$\delta=0.05$	$\delta=0.1$
2	2	2	2	2	2	2	2	2
3	2	2	2	2	2	2	2	2
6	2	2	2	2	3	3	3	3
12	2	2	2	2	3	3	3	3
16	2	2	2	2	3	3	3	3

TABLE VIII
E1: SWISS ROLL ($d_{\text{true}} = 2$), 3-SEED MEAN $\pm$ STD

m	MSE	AU	TwoNN ID	MLE ID
1	1.0160 ± 0.3150	1.0 ± 0.0	1.00 ± 0.00	1.18 ± 0.02
2	0.0290 ± 0.0010	2.0 ± 0.0	1.15 ± 0.05	1.12 ± 0.00
3	0.0010 ± 0.0000	3.0 ± 0.0	1.94 ± 0.06	2.09 ± 0.02
4	0.0010 ± 0.0000	4.0 ± 0.0	1.95 ± 0.06	2.10 ± 0.01

TABLE XI
E3: NSR ANALYSIS (SWISS ROLL, $\sigma = 0.1$)

m	NSR mean	NSR median	Tangent err.	Normal err.
1	6.19	0.586	0.153	0.128
2	1.01	0.562	0.079	0.126
3	1.48	1.010	0.121	0.127
4	1.46	0.988	0.120	0.129

TABLE IX
E1: TORUS ($d_{\text{true}} = 2$), 3-SEED MEAN $\pm$ STD. BOLD ROW ($m = 3$) IS THE EFFECTIVE ELBOW.

m	MSE	AU	TwoNN ID	MLE ID
1	0.36 ± 0.02	1.00 ± 0.00	1.01 ± 0.01	1.12 ± 0.00
2	0.11 ± 0.01	2.00 ± 0.00	1.86 ± 0.04	2.11 ± 0.02
3	0.00 ± 0.00	3.00 ± 0.00	2.02 ± 0.00	2.27 ± 0.03
4	0.00 ± 0.00	4.00 ± 0.00	2.00 ± 0.03	2.28 ± 0.02

TABLE X
E2: DECODER JACOBIAN ANALYSIS

Manifold	m	Sig. SVs	κ	TSA error ($^\circ$)
Swiss Roll	1	1.00	1.00	26.2
	2	2.00	1.51	44.3
	3	3.00	1.75	38.9
	4	3.00	2.03	43.6
Torus	1	1.00	1.00	10.6
	2	2.00	2.30	21.8
	3	3.00	1.50	22.0
	4	3.00	1.47	25.2

normal-direction error for Swiss Roll ($\sigma = 0.1$). Results are shown in Table XI.

At $m = 2 = d_{\text{true}}$, normal-direction error (0.126) $>$ tangent-direction error (0.079), confirming Proposition 2. For

$m > d_{\text{true}}$, surplus dimensions absorb normal-direction noise and selectivity decreases ($\text{NSR} \approx 1.5$).

4) *Trustworthiness vs. Bottleneck Dimension (Experiment E5)*: Table XII and Fig. 3 quantify the m -dependence of Trustworthiness.

Swiss Roll: Trustworthiness jumps from $0.984 \rightarrow 0.999$ at $m = 1 \rightarrow 2$ ($\Delta \approx 0.015$); further increases are < 0.001 . Torus: maximum improvement at $m = 2 \rightarrow 3$ ($\Delta \approx 0.015$), consistent with $d_{\text{emb}} = 3$ (Proposition 4).

C. *Information Bottleneck and Topological Embedding (RQ2, RQ3)*

Question answered (RQ2, RQ3): Does KL-regularized AU saturation reflect the topological minimum embedding dimension? Is decoder isometry ($\kappa \approx 1$) achievable?




Fig. 3. E5: Trustworthiness vs. m (3-seed mean $\pm$ std). Swiss Roll (left): sharp improvement at $m = 2$. Torus (right): sharp improvement at $m = 3 = d_{\text{emb}}$.

TABLE XIII

EA: SWISS ROLL VAE AU vs. m (3-SEED MEAN $\pm$ STD, $\beta = 4$)

m	VAE (AU)	VAE MSE
1	1.0 ± 0.0	1.492 ± 0.257
2	2.0 ± 0.0	0.070 ± 0.012
3	$2.0 \pm 0.0 \leftarrow \text{sat.}$	0.054 ± 0.004
6	2.0 ± 0.0	0.043 ± 0.006
12	2.0 ± 0.0	0.040 ± 0.003
16	2.0 ± 0.0	0.037 ± 0.003

TABLE XIV

EA: TORUS VAE AU vs. m (3-SEED MEAN $\pm$ STD, $\beta = 4$)

m	VAE (AU)	VAE MSE
1	1.0 ± 0.0	0.565 ± 0.029
2	2.0 ± 0.0	0.299 ± 0.009
4	2.0 ± 0.0	0.219 ± 0.005
6	$3.0 \pm 0.0 \leftarrow \text{sat.}$	0.074 ± 0.010
12	3.0 ± 0.0	0.016 ± 0.001
16	3.0 ± 0.0	0.010 ± 0.001

TABLE XV

EB: TOPOLOGICAL AU EXPERIMENT (VAE $\beta = 4$, STANDARDIZED INPUTS). $\checkmark$: $\text{AU}_{\text{sat}} = \text{TOPOLOGICAL EMBEDDING DIM.}$ (MAIN RESULT). $\times$: EXCEPTION (PREPROCESSING-DEPENDENT). Δ : CONDITIONAL.

Manifold	d_{true}	π_1	Topol. d_{emb}	VAE AU_{sat}
Torus T^2	2	$\mathbb{Z}^2$	3	3 $\checkmark$
Möbius Band	2	$\mathbb{Z}$	3	3 $\checkmark$
Swiss Roll	2	$\{1\}$	2	3 $\times$
Circle S^1	1	$\mathbb{Z}$	2	2 ($m \geq 6$) Δ

1) *AU Saturation via Information Bottleneck (Experiment EA)*: Tables XIII and XIV compare $\text{AU}(m)$ for Vanilla AE, AE+L2, and VAE ($\beta = 4$) on Swiss Roll and Torus.

For Swiss Roll, VAE AU maintains $\text{AU} = 2 (= d_{\text{true}})$ for all $m \geq 2$ with zero standard deviation across all 3 seeds—a strong evidence of robustness. For Torus, VAE AU saturates at $3 > d_{\text{true}} = 2$ for $m \geq 6$, corresponding to the topological minimum embedding dimension of 3. Vanilla AE always has $\text{AU} = m$ (confirmed in single-seed experiments), confirming Proposition 3: AU saturation requires KL regularization.

2) *Topological Constraints and AU Saturation (Experiment EB)*: Table XV compares AU_{sat} with the topological minimum embedding dimension for four synthetic manifolds.

Torus and Möbius Band show $\text{AU}_{\text{sat}} = d_{\text{emb}} = 3$ ($\checkmark$: main result). Swiss Roll shows $\text{AU}_{\text{sat}} = 3 \neq d_{\text{emb}} = 2$




Fig. 4. EA: AU vs. m for Swiss Roll (AE vs. VAE). Only VAE exhibits AU saturation at $d_{\text{true}} = 2$.




Fig. 5. EB: VAE AU saturation for four manifolds (standardized inputs). Torus and Möbius Band ($\checkmark$) are the main results. Swiss Roll ($\times$) and Circle S^1 (Δ) are conditional exceptions.

under standardization ($\times$: preprocessing-dependent exception; in Exp. EA without standardization $\text{AU}_{\text{sat}} = 2$ was observed). Circle S^1 requires $m \geq 6$ to observe $\text{AU}_{\text{sat}} = 2$ (Δ : grid-resolution limited).

Note on preprocessing dependence: Experiment EB uses component-wise standardization (mean 0, variance 1), while EA uses the original scale. The Swiss Roll AU_{sat} changing from 2 (EA) to 3 (EB under standardization) shows that AU_{sat} is sensitive to input scale—it is not a preprocessing-invariant topological invariant. The main result (Torus, Möbius Band) should be reproduced under the same standardization protocol.

3) *Isometric Embedding via Decoder Regularization (Experiment EC)*: **Isometry is a prerequisite for reliable ID estimation in latent space:** TwoNN/MLE estimates in latent space are reliable only when the decoder $g: \mathbb{R}^m \rightarrow \mathbb{R}^N$ preserves the distance structure of the data manifold without distortion (isometric embedding), i.e., when $\kappa \approx 1$ (Theorem 2).

TABLE XVI
EC: ISOMETRICAE COMPARISON (SWISS ROLL, $m = 2$, 3-SEED
MEAN $\pm$ STD)

Model	MSE	κ	Trust.
AE	0.028 ± 0.001	1.433 ± 0.182	0.999 ± 0.000
CAE	0.029 ± 0.001	1.907 ± 0.059	0.999 ± 0.000
IsometricAE	0.028 ± 0.001	1.055 ± 0.005	0.999 ± 0.000
DAE	0.030 ± 0.002	1.431 ± 0.144	0.999 ± 0.000




Fig. 6. EC: Condition number κ comparison. IsometricAE achieves the lowest $\kappa \approx 1$.

The IsometricAE loss function is

$$\mathcal{L}_{\text{iso}} = \mathcal{L}_{\text{rec}} + \lambda_{\text{iso}} \cdot \frac{1}{n} \sum_{i=1}^n \|J_g(z_i)^\top J_g(z_i) - I_m\|_F^2 \quad (13)$$

with $\lambda_{\text{iso}} = 0.01$.

Table XVI compares four models on Swiss Roll ($m = 2$, 200 epochs).

IsometricAE achieves the lowest condition number $\kappa = 1.055 \pm 0.005$, confirming Theorem 2. The standard deviation of 0.005 indicates high reproducibility across seeds. CAE shows the highest $\kappa = 1.907 \pm 0.059$ —encoder-side Jacobian regularization unexpectedly distorts the decoder pullback metric. All models achieve Trustworthiness = 0.9990 ± 0.0000 , showing that neighborhood preservation is equivalent across model types at the correct m .

D. Application to MNIST (RQ1 + Novel Finding)

Question (RQ1 on real data + secondary finding): Does the ID estimation framework provide useful design guidance for real image data (MNIST)? What learning dynamics emerge (E6)?

MNIST (10 handwritten digit classes) has a complex manifold structure: each class forms a separate sub-manifold, creating what we term *Manifold Entanglement* [7]—the support $\mathcal{X} = \bigcup_{k=0}^9 \mathcal{M}_k$ cannot be described by a single fundamental group π_1 . This predicts: (1) TwoNN underestimation bias amplification, (2) $\text{AU}_{\text{sat}} > d_{\text{ID}}$ due to class boundary representation, (3) smooth MSE decrease (smooth elbow problem).

TABLE XVII
E4: MNIST AE RESULTS (300 EPOCHS). TWO NN ID AND MLE ID
CONVERGE FOR $m \geq 64$, INDICATING EFFECTIVE ID ≈ 10 –12.

m	MSE	AU	TwoNN ID	MLE ID	CKA
1	0.0437	1	1.02	1.15	0.102
4	0.0292	4	3.86	4.29	0.575
8	0.0191	8	6.08	6.70	0.725
12	0.0148	12	7.59	7.94	0.791
16	0.0121	16	7.96	8.61	0.812
32	0.0087	32	8.84	9.89	0.875
64	0.0062	64	9.54	11.26	0.908
128	0.0061	128	9.84	11.45	0.921
256	0.0057	256	10.02	11.85	0.953

TABLE XVIII
E4: MNIST VAE ($\beta = 4$) RESULTS (300 EPOCHS). AT $m = 12$, AU
GROWTH RATE DROPS SHARPLY ($m^{\text{knee}} = 12$, $\text{AU} = 8$). MNIST IS
GRADUAL-TYPE, SO $m^{\text{knee}} \neq \text{AU}_{\text{sat}}$.

m	MSE	AU	TwoNN ID	MLE ID
1	0.0464	1	1.00	1.11
4	0.0275	4	3.72	4.28
8	0.0186	8	6.36	6.98
12	0.0184	8	6.26	6.97
16	0.0148	11	7.18	8.23
32	0.0121	15	8.25	9.17
64	0.0103	29	9.16	9.87
128	0.0085	26	10.02	10.72
256	0.0072	36	10.14	11.60

TABLE XIX
E6: MNIST TRAINING DYNAMICS ($m = 16$)

Epoch	MSE (test)	Trustworthiness	CKA	TwoNN ID
1	0.0463	0.9263	0.785	4.81
2	0.0358	0.9665	0.832	5.71
5	0.0249	0.9882	0.872	6.67
10	0.0203	0.9912	0.870	6.87
50	0.0158	0.9930	0.853	7.10
100	0.0147	0.9933	0.844	7.52
200	0.0141	0.9929	0.831	7.65

1) *MNIST Bottleneck Sweep (Experiment E4):* We trained AE and VAE ($\beta = 4$) on MNIST (train: 20,000 samples, test: 3,000 samples, 300 epochs) sweeping m from 1 to 256. Results are shown in Tables XVII and XVIII.

From Table XVII, TwoNN ID (10.02) and MLE ID (11.85) converge near $m \geq 64$, suggesting the effective ID of MNIST is in the range of **approximately 10–12**. From Table XVIII, at $m = 12$ the AU growth rate drops sharply ($r(12) = 0/4 = 0.0 < 0.25 \cdot r_{\text{max}}$), identifying $m^{\text{knee}} = 12$. Following the gradual-type design rule (Eq. (6)), the practical bottleneck range is $m^* \approx 12$ –16, consistent with $d_{\text{ID}} \approx 10$ –12.

The high $\text{AU}_{\text{sat}} = 36$ at $m = 256$ has multiple possible interpretations: (1) topological complexity from Manifold Entanglement (class boundary representation), or (2) “capacity leak” (non-essential style variation absorbed at large m). Distinguishing these requires latent traversal studies or persistent homology analysis, which we leave for future work.

2) *MNIST Training Dynamics (Experiment E6):* Table XIX tracks Trustworthiness, CKA, and TwoNN ID over 200 epochs for MNIST ($m = 16$).




Fig. 7. E4: MNIST metric sweeps (300 epochs, 20,000 training samples). Arrow marks $m^{\text{knee}} = 12$ where VAE AU growth rate drops sharply. MSE shows smooth (no clear elbow) decrease—the smooth elbow problem.




Fig. 8. E6: MNIST training dynamics ($m = 16$, 200 epochs). Trustworthiness (global) forms rapidly by Epoch 5; TwoNN ID (local) increases gradually throughout training.

TABLE XX
EX: MNIST AE MULTI-SEED STABILITY (3-SEED MEAN $\pm$ STD, 150 EPOCHS)

m	MSE	AU	TwoNN ID
8	0.018 ± 0.000	8.0 ± 0.0	6.530 ± 0.020
12	0.013 ± 0.000	12.0 ± 0.0	7.870 ± 0.180
16	0.011 ± 0.000	16.0 ± 0.0	8.580 ± 0.200
20	0.010 ± 0.000	20.0 ± 0.0	8.990 ± 0.070
32	0.007 ± 0.000	32.0 ± 0.0	9.720 ± 0.160

Temporal separation of global/local structure formation: Trustworthiness (global neighborhood structure) improves sharply from 0.9263 (Epoch 1) to 0.9882 (Epoch 5). TwoNN ID (local manifold structure) increases gradually from 4.81 to 7.65 throughout training. This temporal separation suggests AEs first capture global data distribution, then refine local manifold geometry.

3) *Multi-Seed Statistical Reliability (Experiment EX):* Tables XX and XXI report 3-seed ($\{42, 123, 456\}$) stability at key m values (150 epochs).

AE MSE standard deviations are within ± 0.0002 and TwoNN ID within ± 0.20 , confirming high reproducibility. VAE AU is perfectly stable at $m = 8$ (std = 0.0) and within ± 1 for $m \geq 12$.

TABLE XXI
EX: MNIST VAE MULTI-SEED STABILITY (3-SEED MEAN $\pm$ STD, 150 EPOCHS)

m	MSE	AU	TwoNN ID
8	0.018 ± 0.000	8.0 ± 0.0	6.400 ± 0.100
12	0.015 ± 0.001	9.7 ± 0.5	7.200 ± 0.110
16	0.013 ± 0.001	12.0 ± 0.8	7.990 ± 0.200
20	0.012 ± 0.001	13.3 ± 0.9	8.330 ± 0.150
32	0.011 ± 0.000	16.7 ± 0.9	9.180 ± 0.280

TABLE XXII
EY: CLASS-RESTRICTED MNIST — AU_{sat} AND TWO NN ID VS. CLASS COUNT (VAE, $\beta = 4$, 75 EPOCHS, VALUES AT $m = 64$)

Config.	K	AU_{sat}	TwoNN ID	MLE ID
1 class (dig. “1”)	1	21	5.10	4.65
2 classes (“0”, “1”)	2	60	5.67	6.28
5 classes (“0”–“4”)	5	50	7.48	8.27
10 classes (all)	10	27	8.89	9.67
(Ref.) E4: 10 cl., 300 ep.	10	36	10.02	11.85

4) *Quantitative Verification of Manifold Entanglement: Class-Restricted MNIST (Experiment EY):* **Objective:** To empirically test the Manifold Entanglement hypothesis introduced in §V-D, we vary the number of MNIST classes $K \in \{1, 2, 5, 10\}$ and examine whether AU_{sat} and TwoNN ID increase systematically with K . $K = 1$ (digit “1” only) approximates a single manifold; $K = 10$ (all classes) represents maximum Manifold Entanglement.

Setup: Samples are drawn uniformly per class ($K = 1$: 2,000 total; $K = 2, 5$: 1,000–500 per class; $K = 10$: 500 per class). VAE ($\beta = 4$, hidden dims [512, 256], 75 epochs) is swept over $m \in \{2, 4, 8, 12, 16, 20, 32, 64\}$ (Experiment EY).

Results and interpretation: TwoNN ID monotonically increases with class count ($5.10 \rightarrow 5.67 \rightarrow 7.48 \rightarrow 8.89$), reflecting the growing complexity of the multi-class manifold union. Comparing the fully converged values—1-class at $AU_{\text{sat}} \approx 7$ (from a 150-epoch preliminary run) versus 10-class at $AU_{\text{sat}} = 36$ (E4)—shows an approximately 5-fold increase, consistent with Manifold Entanglement. Note that the AU_{sat} values from the 75-epoch EY run are preliminary and show non-monotone behavior due to incomplete convergence; a full 150–300 epoch sweep is left as future work.

Conclusion: The monotone increase of TwoNN ID with K quantitatively supports the Manifold Entanglement interpretation of MNIST’s high AU_{sat} : the multi-class manifold structure genuinely increases effective dimensionality, distinguishing it from capacity leak (noise absorption).

VI. DISCUSSION

A. Answer to RQ1: Alignment of m^* with $\hat{d}_{\text{ID}}$

Main conclusion: On synthetic manifolds, m^* aligns well with $\hat{d}_{\text{ID}}$ (Swiss Roll: $m^* = 2 = d_{\text{true}}$; Torus: $m^* = 3 = d_{\text{emb}}$), confirming the ID-based bottleneck design principle.

TwoNN underestimation and mitigation: The underestimation bias at $m = d_{\text{true}}$ (Swiss Roll $m = 2$: ID = 1.15) is a known finite-sample boundary effect [3]; accurate estimates

TABLE XXIII
RQ2 SUMMARY: AU_{sat} VS. TOPOLOGICAL EMBEDDING DIMENSION

Manifold	d_{ID}	π_1	d_{emb}	AU_{sat}
Torus T^2	2	$\mathbb{Z}^2$	3	3 ✓
Möbius Band	2	$\mathbb{Z}$	3	3 ✓
Swiss Roll	2	$\{1\}$	2	3 ×
Circle S^1	1	$\mathbb{Z}$	2	2 △

are obtained at $m > d_{\text{true}}$ ($m = 3$: ID = 1.94). We recommend confirming the stable region across multiple m values.

MNIST application: TwoNN ID = 10.02 and MLE ID = 11.85 from AE ($m = 256$) consistently suggest the effective ID of MNIST is in the range of 10–12 under this setting. The gradual-type AU analysis identifies $m^{\text{knee}} = 12$, consistent with $d_{\text{ID}} \approx 10$ –12.

B. Answer to RQ2: AU_{sat} and Topological Embedding Dimension

Main result (conditional): Under standardized preprocessing (component-wise mean 0, variance 1), Torus and Möbius Band show VAE $AU_{\text{sat}} = 3$, matching the topological minimum embedding dimension.

Why Torus requires $m = 3$: Torus T^2 has $\pi_1(T^2) = \mathbb{Z}^2$; it cannot be continuously embedded in $\mathbb{R}^2$ without self-intersection, requiring at least $\mathbb{R}^3$ [1], [5]. The contractible Gaussian prior support $\mathbb{R}^m$ requires $d_{\text{emb}} > d_{\text{ID}}$ dimensions for embedding nontrivial- π_1 manifolds. This is the theoretical background for why KL regularization drives AU_{sat} toward d_{emb} .

Table XXIII summarizes the RQ2 results.

Preprocessing dependence and the β -data-scale relationship: AU_{sat} depends on the input normalization method and is not a preprocessing-invariant topological invariant. Standardizing the Swiss Roll (mean 0, variance 1) changed AU_{sat} from 2 to 3, demonstrating this sensitivity. This dependence arises from the interplay between the reconstruction loss and the KL term in the VAE objective:

$$\mathcal{L} = \underbrace{\mathcal{L}_{\text{rec}}}_{\text{scale-dependent}} + \beta \underbrace{D_{\text{KL}}(q(z|\mathbf{x})||p(z))}_{\text{scale-invariant}}. \quad (14)$$

The reconstruction loss $\mathcal{L}_{\text{rec}}$ scales with the square of the input variance, while the KL term is scale-invariant. Hence, larger input scale reduces the *effective* KL strength: $\beta_{\text{eff}} \approx \beta/\mathcal{L}_{\text{rec}}$ decreases. When the data scale is large, reconstruction fidelity dominates and the KL pressure on surplus dimensions is weakened, keeping AU_{sat} closer to d_{emb} . Under standardization (scale ≈ 1), the designed β functions as intended, enabling stronger dimension suppression that may activate additional dimensions to encode fine manifold structure (e.g., the longitudinal variation of the Swiss Roll), driving AU_{sat} slightly above d_{emb} . The RQ2 claim $AU_{\text{sat}} \approx d_{\text{emb}}$ should therefore be interpreted as an *empirical correspondence under standardized preprocessing with appropriately balanced β* , not as a strict topological equality. We recommend explicitly reporting the normalization protocol when applying the RQ2 finding.

C. Answer to RQ3: Isometry and Comparison with CAE

Main conclusion: IsometricAE achieves $\kappa = 1.055 \pm 0.005$ (lowest among 4 models), demonstrating that the prerequisite for reliable latent-space ID estimation (isometry) is practically achievable.

CAE comparison: CAE shows the highest $\kappa = 1.907 \pm 0.059$ —encoder-side Jacobian contraction regularization unexpectedly distorts the decoder pullback metric. Isometry requires decoder-side regularization; encoder-side regularization is insufficient.

Trustworthiness limitation: Trustworthiness measures only ordinal preservation of k -nearest-neighbor ranks, not local distance magnitudes. High Trustworthiness does not guarantee $\kappa \approx 1$; the primary metric for RQ3 is κ .

D. Practical Design Guidelines

- 1) **ID estimation via reference AE:** Train an AE with large m (e.g., $m = 256$), apply TwoNN/MLE to its latent space to obtain $\hat{d}_{\text{ID}}$. Confirm stability across multiple m values.
- 2) **VAE sweep and AU pattern observation:** Sweep $\mathcal{M}$ and observe $AU(m)$. For sharp saturation, adopt m^{sat} ; for gradual increase (real data), use m^{knee} as a practical design guideline (note: $AU_{\text{sat}} \neq m^{\text{knee}}$).
- 3) **Finalize with auxiliary criteria:** Within the candidate set, select the minimum m where CKA plateaus (Eqs. (5), (6)). Optionally verify with MSE elbow.

For topologically nontrivial data ($AU_{\text{sat}} > d_{\text{ID}}$ observed), consider setting m^* slightly above $\hat{d}_{\text{ID}}$ (e.g., Torus: $m^* = 3 > d_{\text{ID}} = 2$).

E. Limitations and Future Work

- 1) Experiments are limited to fully-connected AEs; generalization to Conv-AE on CIFAR-10 and natural images requires further investigation.
- 2) EB experiments could be extended to higher-topology manifolds (Klein Bottle, projective spaces).
- 3) The λ_{iso} sensitivity analysis for IsometricAE is not conducted and is left as future work.
- 4) The AU threshold $\delta = 10^{-2}$ has been validated for this study’s data normalization but may need recalibration for significantly different data scales.
- 5) TwoNN underestimation bias at finite samples remains an open problem for quantification.
- 6) The $\tau_{\text{knee}} = 0.25$ threshold has been cross-validated on Fashion-MNIST and Gaussian mixtures (Experiment EZ, Appendix A): MNIST remains robust across $\tau \in [0.1, 0.5]$; Fashion-MNIST shows boundary sensitivity at $\tau = 0.25$ –0.30 (recommend supplementing with the second-order criterion); Gaussian mixture results are dominated by preprocessing-induced posterior collapse. Universality for non-Euclidean data domains and ultra-high-dimensional data remains unverified.
- 7) IsometricAE’s decoder Jacobian computation costs $O(m)$ additional backward passes; scaling to CIFAR-10 ($N = 3072$, $m \approx 32$, $\sim 64\times$ cost) and ImageNet

($\sim 3000\times$) is impractical without approximations (e.g., Hutchinson trace estimation).

- 8) Multi-seed validation of MNIST full-sweep (E4) and learning dynamics (E6) is a future task.
- 9) Manifold Entanglement vs. capacity leak in MNIST's $AU_{\text{sat}} = 36$: Experiment EY (§V-D4) provides initial quantitative evidence via the monotone increase of TwoNN ID with class count K ($5.10 \rightarrow 8.89$). Full 150–300 epoch training across all K values and latent traversal studies are left as future work.
- 10) Confirming Manifold Entanglement requires independent topological analysis via persistent homology [7].
- 11) Quantifying isometry bias in the reference AE latent space for MNIST is unresolved.
- 12) AU_{sat} preprocessing dependence requires a normalization-invariant regularization approach (e.g., Riemannian metric preservation) as a long-term solution.

VII. CONCLUSION

This paper proposed an intrinsic dimension (ID) estimation-based design principle for autoencoder bottleneck dimensions, validated through experiments on synthetic manifolds (Swiss Roll, Torus, Möbius Band) and MNIST. Answers to the three research questions are:

RQ1: TwoNN and MLE estimates $\hat{d}_{\text{ID}}$ align well with m^* on synthetic manifolds, functioning as a strong prior indicator that substantially reduces trial-and-error (with limited accuracy on real data; see §VI-E).

RQ2: On topologically nontrivial manifolds (Torus, Möbius Band), VAE AU_{sat} corresponds to the topological minimum embedding dimension under standardized preprocessing (with exceptions; see §VI-B).

RQ3: IsometricAE achieves decoder condition number $\kappa \approx 1$ (isometry), demonstrating that the prerequisite for reliable latent-space ID estimation is practically achievable.

A secondary finding quantifies the temporal separation of global and local structure formation during MNIST training.

Practical three-step design flow: (1) Apply TwoNN/MLE to a reference AE ($m = 256$) to obtain $\hat{d}_{\text{ID}}$; (2) sweep VAE over grid $\mathcal{M}$ and observe AU pattern (sharp saturation $\Rightarrow m^{\text{sat}}$; gradual increase $\Rightarrow m^{\text{knee}}$); (3) finalize m^* using CKA and MSE as auxiliary criteria.

This work contributes experimental evidence to the growing body of research on ID analysis of deep representations [13], [15] and persistent homology–ID integration [7], [14], from the applied perspective of AE bottleneck design.

REFERENCES

- [1] C. Fefferman, S. Mitter, and H. Narayanan, “Testing the manifold hypothesis,” *Journal of the American Mathematical Society*, vol. 29, no. 4, pp. 983–1049, 2016.
- [2] Y. Bengio, A. Courville, and P. Vincent, “Representation learning: A review and new perspectives,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 35, no. 8, pp. 1798–1828, 2013.
- [3] E. Facco, M. d’Errico, A. Rodriguez, and A. Laio, “Estimating the intrinsic dimension of datasets by a minimal neighborhood information,” *Scientific Reports*, vol. 7, p. 12140, 2017.
- [4] E. Levina and P. J. Bickel, “Maximum likelihood estimation of intrinsic dimension,” in *Advances in Neural Information Processing Systems (NIPS)*, vol. 17, 2005.
- [5] A. Hatcher, *Algebraic Topology*, Cambridge University Press, Cambridge, 2002.
- [6] V. Guillemin and A. Pollack, *Differential Topology*, Prentice-Hall, Englewood Cliffs, NJ, 1974.
- [7] G. Carlsson, “Topology and data,” *Bulletin of the American Mathematical Society*, vol. 46, no. 2, pp. 255–308, 2009.
- [8] D. P. Kingma and M. Welling, “Auto-encoding variational Bayes,” in *International Conference on Learning Representations (ICLR)*, 2014.
- [9] K. Kato, J. Zhou, T. Sasaki, and A. Nakagawa, “Rate-distortion optimization guided autoencoder for isometric embedding in Euclidean latent space,” in *International Conference on Machine Learning (ICML), Proc. Mach. Learn. Res.*, vol. 119, 2020.
- [10] P. Nazari, S. Damrich, and F. A. Hamprecht, “Geometric autoencoders – what you see is what you decode,” in *International Conference on Machine Learning (ICML), Proc. Mach. Learn. Res.*, vol. 202, 2023.
- [11] C. Ceruti, S. Bassis, A. Rozza, G. Lombardi, E. Casiraghi, and P. Campadelli, “DANCo: An intrinsic dimensionality estimator exploiting angle and norm concentration,” *Pattern Recognition*, vol. 47, no. 8, pp. 2569–2581, 2014.
- [12] P. Pope, C. Zhu, A. Abdelkader, M. Goldblum, and T. Goldstein, “The intrinsic dimension of images and its impact on learning,” in *International Conference on Learning Representations (ICLR)*, 2021.
- [13] A. Ansuini, A. Laio, J. H. Macke, and D. Zoccolan, “Intrinsic dimension of data representations in deep neural networks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [14] T. Birdal, A. Lou, L. J. Guibas, and U. Simsekli, “Intrinsic dimension, persistent homology and generalization in neural networks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021.
- [15] L. Valeriani, D. Wood, G. Catania, S. Gherardini, A. Laio, and A. Ansuini, “The geometry of hidden representations of large transformer models,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2023.
- [16] I. Higgins, L. Matthey, A. Pal, C. Burgess, X. Glorot, M. Botvinick, S. Mohamed, and A. Lerchner, “ β -VAE: Learning basic visual concepts with a constrained variational framework,” in *International Conference on Learning Representations (ICLR)*, 2017.
- [17] J. Lucas, G. Tucker, R. B. Grosse, and M. Norouzi, “Don’t blame the ELBO! A linear VAE perspective on posterior collapse,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [18] S. Rifai, P. Vincent, X. Muller, X. Glorot, and Y. Bengio, “Contractive auto-encoders: Explicit invariance during feature extraction,” in *International Conference on Machine Learning (ICML)*, pp. 833–840, 2011.
- [19] G. Alain and Y. Bengio, “What regularized auto-encoders learn from the data-generating distribution,” *Journal of Machine Learning Research (JMLR)*, vol. 15, pp. 3743–3773, 2014.
- [20] S. Kornblith, M. Norouzi, H. Lee, and G. Hinton, “Similarity of neural network representations revisited,” in *International Conference on Machine Learning (ICML), Proc. Mach. Learn. Res.*, vol. 97, pp. 3519–3529, 2019.
- [21] N. Tishby, F. C. Pereira, and W. Bialek, “The information bottleneck method,” in *37th Annual Allerton Conference*, pp. 368–377, 2000.
- [22] Y. Burda, R. B. Grosse, and R. Salakhutdinov, “Importance weighted autoencoders,” in *International Conference on Learning Representations (ICLR)*, 2016.
- [23] L. van der Maaten and G. Hinton, “Visualizing data using t-SNE,” *Journal of Machine Learning Research (JMLR)*, vol. 9, pp. 2579–2605, 2008.

APPENDIX A

EXPONENTIAL SATURATION MODEL AND SENSITIVITY ANALYSIS FOR τ_{knee}

Information-Theoretic Motivation

If $AU(m) \approx A(1 - e^{-\alpha m})$, then the growth rate $r(m) \propto e^{-\alpha m}$ decays exponentially. The natural inflection of exponential decay occurs around $r(m) \approx r_{\text{max}}/e \approx 0.37 \cdot r_{\text{max}}$, and $\tau_{\text{knee}} = 0.25$ is positioned as a conservative lower bound of this theoretical inflection region $[0.25, e^{-1}]$. Thus $r(m) < 0.25 \cdot r_{\text{max}}$ means “information utilization efficiency

TABLE XXIV
EZ: τ_{knee} CROSS-DATASET ROBUSTNESS

Dataset	AU _{sat}	TwoNN	$\tau=0.1$	$\tau=0.2$	$\tau=0.25$	$\tau=0.3$	$\tau=0.4$	$\tau=0.5$
Fashion-MNIST	10	7.00	20	20	20	8	8	8
GMix (5-comp)	1	4.27	2	2	2	2	2	2
GMix (10-comp)	3	3.02	2	2	2	2	2	2
MNIST (ref.)	36	10.02	12	12	12	12	12	12

below 25% of peak,” indicating effective dimension usage has essentially ceased.

Sensitivity Analysis

For MNIST VAE, $r(12) = 0.0$ and $r(16) = 0.75$, so any $\tau_{\text{knee}} \in (0, 0.75)$ selects $m^{\text{knee}} = 12$ (robust over a wide interval of length 0.75). Across all datasets (Swiss Roll, Torus, MNIST), we confirmed that m^{knee} identification is unchanged for any $\tau_{\text{knee}} \in [0.1, 0.5]$.

Multi-seed experiments (EX) with seeds $\{42, 123, 456\}$ confirm that $r(12)$ drops sharply in all seeds, yielding $m^{\text{knee}} = 12$ consistently (EX representative values: $m = 12$ AU $= 7.7 \pm 0.6$, $r(12) = 0.00 \pm 0.00$, $r(16) = 0.8 \pm 0.2$).

Cross-Dataset Validation of τ_{knee} Robustness (Experiment EZ)

To assess the universality of $\tau_{\text{knee}} = 0.25$ beyond MNIST, we evaluate Fashion-MNIST and synthetic Gaussian mixtures (5- and 10-component, $d_{\text{true}} = 5$, $D = 50$) using VAE ($\beta = 4$), and compare m^{knee} across $\tau \in \{0.1, 0.2, 0.25, 0.3, 0.4, 0.5\}$ (Experiment EZ).

Fashion-MNIST: AU increases gradually from 1 to 10 over $m = 1$ to 64 (gradual-type, like MNIST; TwoNN ID ≈ 7.0). The normalized rate $r(m) = 0.25$ (exactly equal to τ_{knee}) at $m = 8, 12, 16$, causing the strict “ $<$ ” condition to skip to $m^{\text{knee}} = 20$ for $\tau = 0.25$ while $\tau = 0.30$ yields $m^{\text{knee}} = 8$. This boundary sensitivity — arising from the rate being exactly at the threshold — suggests supplementing with a data-driven second-order difference criterion (see §VI-E) when $r(m)$ values cluster near τ_{knee} . The qualitative conclusion (“ $m = 8$ – 20 is the slow-growth region”) is consistent across all $\tau \in [0.1, 0.5]$.

Gaussian mixtures: The $[0, 1]$ normalization of the Gaussian mixture data dramatically amplifies the effective KL strength (see preprocessing dependence discussion, §VI-B), causing posterior collapse (AU_{sat} = 1–3). The $m^{\text{knee}} = 2$ identification is stable across all τ , though it does not reflect $d_{\text{true}} = 5$ due to the preprocessing-induced collapse. Applying component-wise standardization instead of $[0, 1]$ normalization would recover more meaningful AU_{sat} values.

Summary: $\tau_{\text{knee}} = 0.25$ is fully robust for MNIST across $[0.1, 0.5]$. For Fashion-MNIST, boundary sensitivity arises at $\tau = 0.25$ – 0.30 ; the supplementary second-order criterion resolves ambiguity. Gaussian mixture results are dominated by preprocessing-induced posterior collapse rather than τ sensitivity.