

Knowledge Extraction from Autoencoder Latent Spaces: Diagnosing Effective Dimensionality via Active-Unit Dynamics and Intrinsic-Dimension Estimates

Chika Obata ¹, Mizuki Dai ¹, and Kenya Jin'no ^{1,*}

¹ Department of Intelligent Systems, Faculty of Information Engineering, Tokyo City University, 1-28-1 Tamazutsumi, Setagaya-ku, Tokyo 158-8557, Japan; {obata, dai}@tcu.ac.jp

* Correspondence: kjinno@tcu.ac.jp; Tel.: +81-3-5707-0104

Knowledge Extraction from Autoencoder Latent Spaces: Diagnosing Effective Dimensionality via Active-Unit Dynamics and Intrinsic-Dimension Estimates

Chika Obata ¹, Mizuki Dai ¹, and Kenya Jin'no ^{1,*}

¹ Department of Intelligent Systems, Faculty of Information Engineering, Tokyo City University, 1-28-1 Tamazutsumi, Setagaya-ku, Tokyo 158-8557, Japan; {obata, dai}@tcu.ac.jp

* Correspondence: kjinno@tcu.ac.jp; Tel.: +81-3-5707-0104

Abstract: Determining how many latent coordinates an autoencoder actually uses is difficult because reconstruction error, intrinsic-dimension estimates, and variational activity reflect different aspects of a learned representation. This paper proposes a multi-indicator diagnostic framework for extracting actionable knowledge from autoencoder latent spaces. The framework combines active-unit (AU) dynamics, TwoNN/MLE intrinsic-dimension (ID) estimates, reconstruction-error elbows, and PCA baselines, while explicitly separating their diagnostic roles. AU counts are treated as direct observations of coordinate usage in VAEs; TwoNN/MLE estimates as numerically stable auxiliary references subject to geometric bias; and the AU slowdown point, m^{knee} , as an exploratory, grid- and seed-dependent change point rather than a unique optimal bottleneck size.

Under a linear-Gaussian β -VAE rate–distortion surrogate, we derive the activation condition $\lambda_j > \beta\tau^2$, providing an interpretable water-filling view of KL-induced dimension suppression. Experiments on fully connected VAEs trained on MNIST show that AU counts and TwoNN estimates are stable across random seeds, whereas sparse-grid m^{knee} is unreliable (43 ± 15). A dense grid around the estimated intrinsic dimension stabilizes m^{knee} at 10.3 ± 0.9 , consistent with TwoNN near 10. Additional experiments on Conv-VAEs, CIFAR-10, and SVHN show that stable identification is not always possible and clarify the boundary conditions under which $\text{AU}/m^{\text{knee}}$ diagnostics should be interpreted as inconclusive. These results support the proposed framework as a reproducible diagnostic protocol for latent-dimensionality analysis and model-design support, rather than a universal dimension estimator.

Keywords: variational autoencoder; effective latent dimensionality; active units; intrinsic dimension; TwoNN; diagnostic framework; representation learning; knowledge extraction; posterior collapse

1. Introduction

1.1. Background

Autoencoders and variational autoencoders (VAEs) [11] learn compact representations of high-dimensional data by forcing information through a bottleneck of dimension m . A key design question is: *how many of the m latent coordinates does the model actually use?* This “effective latent dimensionality” governs reconstruction fidelity, disentanglement, and posterior collapse [19].

In the context of machine learning and knowledge extraction, latent representations underpin model selection, compression, downstream prediction, clustering, and visualization [7]. However, the nominal bottleneck dimension m does not necessarily coincide with the number of coordinates the model actually exploits. A diagnostic protocol that distinguishes coordinate usage, geometric intrinsic dimension, and reconstruction-side behavior is therefore directly useful for interpreting learned representations and selecting latent-space sizes in a reproducible manner. From the viewpoint of knowledge extraction, the objective of this study is not merely to estimate a single latent dimensionality value. Rather, we aim to extract model-design knowledge from trained autoencoder representations. The proposed diagnostics answer three practical questions: which latent coordinates are actually used by the model; whether the geometry of the learned representation is consistent with intrinsic-dimension estimates; and whether increasing the nominal bottleneck dimension leads to meaningful representational gains. This knowledge is useful for selecting latent-space sizes, detecting over-parameterized or under-regularized representations, and identifying cases in which dimensionality diagnostics should be regarded as inconclusive. In this sense, the proposed framework provides a knowledge-extraction protocol for representation learning, rather than a standalone intrinsic-dimension estimator.

Existing work measures individual aspects of latent geometry—AU counts [18, 19], ID estimates [8, 9], isometric regularization [12]—but a systematic framework that combines these indicators, clarifies their distinct roles, and defines the conditions under which each is reliable has not been established.

1.2. Problem Statement and Contributions

The central research question is:

How can we diagnose, without over-claiming, how many dimensions an AE/VAE effectively exploits,

and where does such diagnosis break down?

The contributions of this paper are threefold.

1. **Multi-indicator diagnostic framework.** We propose a framework that explicitly assigns distinct roles to four indicator classes: (i) AU values as direct usage observations; (ii) TwoNN/MLE ID as numerically stable auxiliary references subject to geometric bias; (iii) m^{knee} as an exploratory change point; and (iv) MSE elbow and PCA as reconstruction-side baselines. The main contribution is not m^{knee} itself but the role separation, which prevents users from over-relying on any single indicator. MNIST three-seed experiments confirm that AU values ($\sigma \leq 0.5$) and TwoNN ID (10.03 ± 0.12) are reproducible; the main stabilized result for m^{knee} is 10.3 ± 0.9 under dense-grid, independent-initialization, multi-seed conditions.
2. **Linear-Gaussian rate–distortion analysis.** Under a linear-Gaussian β -VAE rate–distortion approximation, we derive the activation condition $\lambda_j > \beta\tau^2$ and show that AU saturation follows a water-filling principle. This analytical result is an *interpretive approximation model* for the nonlinear setting, not a general theorem for arbitrary VAEs.
3. **Empirical characterization of reliability and limits.** Dense-grid/multi-seed experiments show that sparse-grid m^{knee} is unreliable ($\sigma = 15$), whereas dense-grid evaluation stabilizes it near the TwoNN estimate. Conv-VAE and natural-image experiments identify boundary conditions under which AU/ m^{knee} diagnostics become unreliable or require stronger regularization and architectural control.

1.3. Paper Organization

Section 2 reviews related work. Section 3 defines the diagnostic framework. Section 4 provides theoretical background. Section 5 reports experiments. Section 6 discusses findings. Section 7 concludes. Detailed tables and proofs are in the Appendix.

2. Related Work

Representation learning and autoencoders. Bengio et al. [7] provide a comprehensive review of representation learning, including the manifold hypothesis [6] underpinning AEs. β -VAE [18] increases KL weight to encourage disentanglement; Lucas et al. [19] analyse posterior collapse in the linear-VAE limit. Contractive AEs [20] and denoising AEs [21] offer complementary regularization strategies.

Intrinsic-dimension estimation. TwoNN [8] and MLE [9] estimate local intrinsic dimension from nearest-neighbour distances; DANC0 [14] offers higher accuracy at greater computational cost. Pope et al. [15] apply TwoNN to natural images, reporting $d_{\text{ID}} \approx 26$ for CIFAR-10 at 32×32 . Ansuini et al. [16] track ID across deep-network layers; Valeriani et al. [17] extend this to large transformer models.

Geometric autoencoders. Nazari et al. [13] and Kato et al. [12] propose isometric regularization of the decoder Jacobian. We use IsometricAE as a controlled testbed for geometric bias, not as the main tool.

Positioning. The present work differs from the above in that it targets the *design stage*—systematically studying how AU dynamics and ID estimation jointly characterize effective dimensionality, and quantifying when each indicator is reliable.

3. Diagnostic Framework

3.1. Problem Setup

Let an AE/VAE have encoder $f : \mathbb{R}^N \rightarrow \mathbb{R}^m$ and decoder $g : \mathbb{R}^m \rightarrow \mathbb{R}^N$. For a VAE, the encoder outputs a posterior $q_\phi(z|x) = \mathcal{N}(\mu(x), \sigma^2(x))$; KL regularization (β -weighted) acts as an information bottleneck.

Definition 3.1 (Active Units [19]).

$$\text{AU}(m) = \sum_{k=1}^m \mathbf{1}[\text{Var}_x[\mu_k(x)] > \delta], \quad \delta = 10^{-2} \quad (1)$$

where $\mu_k(x)$ is the k -th posterior mean. AU is robust to the choice of $\delta \in [10^{-3}, 10^{-1}]$ (Appendix E.1).

Definition 3.2 (AU Slowdown Point). Let $s(m_i) = (\text{AU}(m_i) - \text{AU}(m_{i-1})) / (m_i - m_{i-1})$ be the raw AU increase rate, and $\rho(m_i) = s(m_i) / \max_j s(m_j) \in [0, 1]$ the normalised rate. The slowdown point is

$$m^{\text{knee}} = \min\{m_i \in \mathcal{M} : \rho(m_i) < \tau_{\text{knee}}\}, \quad \tau_{\text{knee}} = 0.25. \quad (2)$$

Role separation. Table 1 summarises the roles.

Table 1. Classification and empirical reliability of diagnostic indicators. *Robust* means low inter-seed variance under tested conditions, not unbiasedness toward the true ID. *Exploratory* means useful for hypothesis generation but not a unique estimate of effective dimensionality.

Indicator	Category	Empirical reliability	Primary use
AU values	Stable across seeds [‡]	$\sigma \leq 0.5$ (MNIST 3-seed)	Direct observation of active coords (not unbiased toward true ID)
TwoNN/MLE ID	Numerically stable auxiliary [†]	$\sigma \leq 0.12$ ($m=64$)	Reference ID estimate (stable, not unbiased)
MSE elbow	Auxiliary	Qualitative only	Reconstruction-side check
PCA cum. variance	Auxiliary	Linear assumption	Linear baseline
m^{knee}	Exploratory	$\sigma = 15$ sparse; 0.9 dense	Rough diagnostic heuristic

[†]Subject to isometry bias; numerically stable but not unbiased toward true d_{ID} . [‡]“Stable across seeds” means low inter-seed variance under tested conditions, not unbiasedness. m^{knee} is not a standalone estimator; informative only under dense-grid, independent-initialization, multi-seed evaluation.

3.2. Two-Stage Diagnostic Pipeline

Stage 1 (ID estimation, auxiliary). For synthetic data with known d_{true} , TwoNN/MLE is applied directly. For real data (e.g. MNIST), a large reference AE ($m_{\text{ref}} = 256$) is trained; its latent representation is fed to TwoNN/MLE to obtain $\hat{d}_{\text{ID}}$. *Limitation:* this reference AE is not isometric ($\kappa \approx 10^9$); $\hat{d}_{\text{ID}}$ is an approximation subject to geometric bias.

Stage 2 (AU Dynamics, primary). A VAE ($\beta = 4$) is trained on a grid $\mathcal{M} = \{1, 2, 4, 8, 12, 16, 20, 32, 64, 128, 256\}$. For *sharp-saturation* data (synthetic manifolds), AU_{sat} is read directly. For *gradual-increase* data (real data), m^{knee} is computed from Eq. (2) and cross-checked against $\hat{d}_{\text{ID}}$ from Stage 1.

4. Theoretical Background

4.1. Linear-Gaussian Rate–Distortion Analysis

Analytical Result 4.1 (AU Activation Condition in Linear-Gaussian Rate–Distortion Surrogate). *Consider a β -VAE with linear encoder/decoder and Gaussian prior, analyzed under a linear-Gaussian rate–distortion surrogate approximation (details in Appendix B.2). Under this approximation, the water-filling condition implies:*

1. **(Activation threshold)** *The j -th eigen-direction becomes active iff $\lambda_j > \beta\tau^2$, where λ_j is the j -th data-covariance eigenvalue, τ^2 is the decoder noise variance.*
2. **(Upper bound)** *The optimal number of active eigen-directions satisfies $A^*(m) \leq \min(m, d^*(\beta))$, where $d^*(\beta) := |\{j : \lambda_j > \beta\tau^2\}|$.*
3. **(Monotone saturation of optimal solution)** *$A^*(m)$ is non-decreasing in m and saturates at $d^*(\beta)$ for $m \geq d^*(\beta)$.*

Remark 4.2 (Empirical AU is non-monotone). *Analytical Result 4.1(3) concerns $A^*(m)$, the global-optimum active eigen-direction count. In practice, each bottleneck size m is trained independently with stochastic optimization; the resulting empirical $AU(m)$ can be non-monotone due to finite-sample effects and initialization dependence (e.g. $AU=29$ at $m=64$ but $AU=26$ at $m=128$, Table 8). Therefore, m^{knee} should be treated as an exploratory diagnostic indicator, not a precise estimate of $d^*(\beta)$.*

Remark 4.3 (Vanilla AE corresponds to $\beta=0$). *Without KL regularization, the threshold $\beta\tau^2$ vanishes, so all directions with positive variance are active. Thus the pseudo-AU (coordinate-usage count) of a standard AE approximates m .*

4.2. Grid-Resolution Dependence of m^{knee}

Proposition 4.4 (Grid Resolution Dependence). *Under an idealised monotone-saturation model where $AU(m)$ changes like a step function at $d^*(\beta)$ without noise, the detection error satisfies $|m^{\text{knee}}(\mathcal{G}) - d^*(\beta)| \leq \max_i(m_{i+1} - m_i)$.*

Remark 4.5. *This bound applies to the idealised model, not real non-linear VAEs whose AU curves fluctuate non-monotonically. The practical implication is qualitative: finer grids near the estimated $\hat{d}_{\text{ID}}$ reduce the variance of m^{knee} , as confirmed in Section 5.4.*

4.3. Pullback Metric and Isometry

Proposition 4.6 (Pullback Metric and Local Isometry). *The decoder $g : \mathbb{R}^m \rightarrow \mathbb{R}^N$ is locally isometric at z iff $G(z) = J_g(z)^\top J_g(z) = I_m$ (all singular values equal 1). The condition number $\kappa = \sigma_{\max}/\sigma_{\min}$ measures anisotropy: $\kappa=1$ means all singular values are equal but not necessarily 1. Therefore κ is an anisotropy indicator, not a sufficient condition for isometry.*

4.4. Design Principle: Bottleneck Margin is a Major TwoNN Stabilizing Factor

Experiments in Section 5.5 support the following design principle.

Design Principle 1. Under the tested conditions, the numerical stability of TwoNN estimates is explained more consistently by bottleneck margin ($m \gg \hat{d}_{\text{ID}}$) than by isometric regularization alone. TwoNN underestimates when $m \approx d_{\text{true}}$; it converges once $m \geq 5\hat{d}_{\text{ID}}$ approximately.

This is an empirically supported guideline verified primarily on MNIST multi-seed experiments and qualitatively checked on Fashion-MNIST. It should not be interpreted as a claim that isometric regularization is irrelevant; rather, bottleneck margin was the dominant stabilizing factor in the tested settings.

5. Experiments

5.1. Experimental Setup

All models use Adam ($\text{lr} = 10^{-3}$), batch size 128, ReLU activations, Sigmoid output, and PyTorch 2.0.

Synthetic manifolds. Hidden layers [64, 32], 200–300 epochs, $\beta = 4.0$, $n = 5000$.

MNIST / Fashion-MNIST. Input $N = 784$, hidden [512, 256], $\beta = 4.0$, grid $\mathcal{M} = \{1, 2, 4, 8, 12, 16, 20, 32, 64, 128, 256\}$.

MNIST: 300 epochs, $n_{\text{train}} = 20\,000$, $n_{\text{test}} = 3\,000$. Fashion-MNIST: 100 epochs, $n_{\text{train}} = 10\,000$.

ID estimation: TwoNN ($k = 2$), MLE ($k = 10$), AU threshold $\delta = 10^{-2}$. Three random seeds {42, 123, 777} for main reproducibility experiments. All experiments in the multi-seed study use independent per- (m, seed) initialization.

5.2. Synthetic Manifolds: Principle Verification

5.2.1. MSE Elbow and Manifold Dimension

On Swiss Roll ($d_{\text{true}} = 2$) and Torus ($d_{\text{true}} = 2$, $d_{\text{emb}} = 3$), the reconstruction MSE decreases substantially from $m = 1$ to $m = 2$, and further at $m = 3$ (Table 2). For Swiss Roll, the decrease from $m = 1$ to $m = 2$ is consistent with the idea that a two-dimensional bottleneck is sufficient to represent this manifold in practice. For Torus, the largest decrease occurs at $m = 3$, consistent with the three-dimensional embedding requirement. However, the anomalous increase at $m = 4$ for Swiss Roll (0.041 ± 0.015 , higher than at $m = 2$) illustrates that MSE elbows are affected by optimization variance, network capacity, and coordinate representation, not only by intrinsic dimensionality. In addition, simple connectedness alone does not guarantee that the embedding dimension equals the intrinsic dimension. Therefore, the MSE elbow should be treated only as a reconstruction-side diagnostic indicator, cross-checked with AU and TwoNN, rather than as a direct estimate of intrinsic dimension (Appendix B).

Table 2. MSE of AE on synthetic manifolds (mean $\pm$ std, 3 seeds). Swiss Roll elbow at $m = d_{\text{true}} = 2$; Torus elbow at $m = d_{\text{emb}} = 3$. The rise at $m = 4$ for Swiss Roll (std= 0.015, vs. 0.002 at $m = 2$) reflects optimization variance when the extra dimension is unused: higher std indicates sensitivity to random initialization.

m	Swiss Roll MSE	Torus MSE	Note
1	1.016 ± 0.057	0.388 ± 0.017	under-parameterised
2	0.029 ± 0.002	0.085 ± 0.004	elbow (Swiss Roll)
3	0.0008 ± 0.0001	0.0005 ± 0.0000	elbow (Torus)
4	0.041 ± 0.015	0.019 ± 0.002	high std: optimization variance

5.2.2. AU Saturation on Synthetic Data

VAE ($\beta = 4$) saturates at $\text{AU}_{\text{sat}} = 2$ (Swiss Roll, $= d_{\text{true}}$) and $\text{AU}_{\text{sat}} = 3$ (Torus, $= d_{\text{emb}} > d_{\text{true}}$), both with zero inter-seed variance, consistent with Analytical Result 4.1(qualitative). For standard AE, pseudo-AU= m (all coordinates active), consistent with the $\beta = 0$ remark of Analytical Result 4.1. Detailed topological-manifold results (Torus, Möbius Band) are in Appendix C.

5.3. MNIST: Multi-seed Reproducibility Study

5.3.1. ID Estimation via Reference AE

A standard AE trained on MNIST (300 epochs) with $m = 256$ yields TwoNN ID ≈ 10.02 , MLE ID ≈ 11.85 . Multiple m values confirm convergence above $m \geq 64$: TwoNN ID = 9.54 at $m = 64$, stabilising to 10.02 at $m = 256$. This places MNIST's intrinsic dimension at $\hat{d}_{\text{ID}} \approx 10$ –12 (Table 3).

Table 3. MNIST AE results (300 epochs, seed 42).

m	MSE	TwoNN ID	MLE ID
1	0.0437	1.02	1.15
32	0.0087	8.84	9.89
64	0.0062	9.54	11.26
256	0.0057	10.02	11.85

5.3.2. VAE AU Dynamics: Gradual Growth Observed

MNIST VAE ($\beta = 4$) shows gradual, non-monotone AU growth as m increases, consistent with non-linear stochastic optimization (e.g., AU = 29 at $m = 64$ but AU = 26 at $m = 128$). This gradual growth is the key reason why sparse-grid m^{knee} is unreliable: the AU curve does not have the sharp step assumed in the idealised model (Proposition 4.4). A single-seed preliminary run illustrating this pattern is provided in Appendix A (Table 8). **The $m^{\text{knee}} = 12$ value observed in that run is not used as a main result**; the main stabilized result is $m^{\text{knee}} = 10.3 \pm 0.9$ from dense-grid, multi-seed evaluation (Section 5.4).

5.3.3. Multi-seed Experiments: AU and TwoNN are Stable, m^{knee} is Not

Table 4 reports MNIST VAE results with sparse grid $\mathcal{M} = \{4, 8, 12, 16, 20, 32, 64\}$ and three seeds (300 epochs).

Table 4. MNIST VAE multi-seed study: Exp-Mnist-Multi-300 (mean $\pm$ std, 3 seeds, sparse grid). AU and TwoNN ID are reproducible; m^{knee} varies across seeds.

m	AU (mean $\pm$ std)	TwoNN ID (mean $\pm$ std)	Note
4	4.0 ± 0.0	3.98 ± 0.14	all active
12	10.0 ± 0.0	7.41 ± 0.06	
32	15.7 ± 0.5	9.07 ± 0.10	
64	22.3 ± 0.5	10.03 ± 0.12	converges to $\hat{d}_{\text{ID}} \approx 10$
Sparse-grid m^{knee} ($\tau = 0.25$): seed 42=64, 123=32, 777=32; mean 43 ± 15			

Key findings. (1) AU values and TwoNN ID are highly reproducible across seeds ($\sigma \leq 0.5$ and $\sigma \leq 0.12$, respectively). (2) Sparse-grid m^{knee} varies from 32 to 64 ($\sigma = 15$), confirming that it is an unreliable diagnostic reference at coarse grids.

5.4. Dense-Grid Stabilization of m^{knee}

To isolate the effect of grid density, we designed Exp-Mnist-Dense-300 with $\mathcal{M} = \{8, 9, 10, 11, 12, 13, 14, 15, 16, 18, 20, 24, 28, 32, 40\}$ independent per- (m, seed) initialization, 300 epochs, 3 seeds.

The dense-grid result $m^{\text{knee}} = 10.3 \pm 0.9$ is consistent with TwoNN ID (≈ 10) and represents the main stabilized m^{knee} result of this paper. Even with dense grids the AU curve fluctuates non-monotonically (AU: 8 $\rightarrow$ 9 $\rightarrow$ 10 $\rightarrow$ 10 $\rightarrow$ 10 $\rightarrow$ 10 $\rightarrow$ 12 $\rightarrow$ 11 $\rightarrow$ 14 $\rightarrow$...), so m^{knee} should still be treated as an exploratory heuristic, not a

Table 5. Dense-grid vs. sparse-grid m^{knee} comparison (MNIST FC-VAE, $\beta = 4$). Grid density is the dominant stabilization factor.

Condition	Max. grid interval	seed 42	seed 123	seed 777	Mean $\pm$ std
Sparse, 300ep (Exp-Mnist-Multi-300)	32	64	32	32	43 ± 15
Dense, 300ep (Exp-Mnist-Dense-300)	8	11	9	11	10.3 ± 0.9

precise estimator. The fixed- τ method outperforms the second-difference method for dense grids ($\sigma = 0.9$ vs. 18.3 ± 6.8 , Appendix C).

Figure 1 visualises the sparse-grid instability and dense-grid stabilization.

Figure 2: MNIST VAE AU dynamics — sparse vs. dense grid




Figure 1. MNIST VAE AU dynamics. **Left:** Sparse grid ($m \in \{4, 8, \dots, 64\}$, 3 seeds, 300 epochs)— m^{knee} varies between 32 and 64 ($\sigma = 15$). **Right:** Dense grid ($m \in \{8, 9, \dots, 64\}$, independent initialization, 3 seeds, 300 epochs)— $m^{\text{knee}} = 10.3 \pm 0.9$, consistent with TwoNN ID (≈ 10 , dashed red). The non-monotone AU fluctuations visible in both panels confirm that m^{knee} is an exploratory heuristic requiring controlled evaluation conditions.

5.5. TwoNN Stability: Bottleneck Margin as a Major Factor

Three complementary experiments support Design Principle 1 (Figure 2).

5.5.1. Swiss Roll: Isometry vs. Margin

For standard AE ($\kappa = 1.79$) and IsometricAE ($\kappa = 1.06$), TwoNN ≈ 1.55 at $m = 2 = d_{\text{true}}$ for both—the bottleneck constraint drives underestimation, not κ difference. At $m = 4 > d_{\text{true}}$, TwoNN recovers to ≈ 2.03 for standard AE despite $\kappa = 2.16$ (Appendix C, Table 9).

5.5.2. MNIST AE m -Sweep

TwoNN stabilizes at $m \geq 48 \approx 5\hat{d}_{\text{ID}}$ (≈ 10.16) and remains stable even when κ jumps from 4.25 to 37.0 between $m = 64$ and $m = 128$ (TwoNN: $10.17 \rightarrow 10.98$, $\Delta < 0.82$; color-coded in Figure 2).

5.5.3. Four-Model Comparison on MNIST (Exp-Real-Geom)

AE, IsometricAE, CAE, and DAE are swept over $m \in \{d, 2d, 3d, 5d, 10d\}$ ($d = \hat{d}_{\text{ID}} = 10$). All four models converge at $m = 5d$ – $10d$; inter-model TwoNN difference at the same m is < 0.3 , far smaller than the gain from




Figure 2. TwoNN ID vs. bottleneck dimension m for MNIST AE (color = $\log_{10} \kappa$, anisotropy indicator). TwoNN stabilizes above $m \approx 5\hat{d}_{ID}$ (vertical dotted line), while κ increases substantially beyond $m = 64$ without affecting stability. This supports Design Principle 1: bottleneck margin, not isometry, is the major stabilizing factor in the tested settings.

increasing m (~ 2.5 points). These results strongly support Design Principle 1 (see Table 6 in main text and Appendix C for Fashion-MNIST).

Table 6. Exp-Real-Geom: MNIST four-model comparison (AE, IsometricAE, CAE, DAE) (mean $\pm$ std, 3 seeds; $\hat{d}_{ID} = 10$, representative m values). TwoNN stabilization is driven by $m/\hat{d}_{ID}$ across all four models; inter-model difference at the same m is ≤ 0.3 .

Model	m	TwoNN ID	MLE ID	κ	Trust.	kNN ovlp.
AE	10 ($= d$)	7.3 ± 0.4	8.1 ± 0.3	85 ± 22	0.961	0.73
	30 ($= 3d$)	9.8 ± 0.2	10.5 ± 0.2	1200 ± 400	0.981	0.85
	100 ($= 10d$)	10.0 ± 0.1	10.8 ± 0.1	4.1×10^7	0.987	0.88
IsometricAE	10	7.5 ± 0.4	8.3 ± 0.3	3.2 ± 0.4	0.963	0.74
	30	9.9 ± 0.1	10.6 ± 0.1	5.2 ± 0.8	0.982	0.86
	100	10.1 ± 0.1	10.8 ± 0.1	7.8 ± 1.4	0.988	0.89
CAE	10	7.2 ± 0.5	8.0 ± 0.4	12.1 ± 3.2	0.958	0.72
	30	9.7 ± 0.2	10.4 ± 0.2	95 ± 31	0.980	0.84
	100	10.0 ± 0.1	10.7 ± 0.1	3200 ± 1100	0.986	0.87
DAE	10	7.6 ± 0.4	8.4 ± 0.4	18.5 ± 5.4	0.964	0.75
	30	9.9 ± 0.1	10.5 ± 0.2	183 ± 62	0.982	0.85
	100	10.0 ± 0.1	10.8 ± 0.1	6300 ± 2100	0.987	0.88

5.6. Fashion-MNIST: Preliminary External Check

Fashion-MNIST is used only as a preliminary external check, not as evidence of seed-level robustness. Single-seed (seed 42) results: TwoNN ID converges to ≈ 9.4 at $m \geq 32$, suggesting $\hat{d}_{ID} \approx 9\text{--}10$ (slightly lower than MNIST). VAE ($\beta = 4$) shows a sharp $\rho(m)$ drop around $m = 12$ in this single run, but multi-seed reproducibility is not assessed. Fashion-MNIST 3-seed experiments are left for future work.

5.7. Conv-VAE: Applicability Boundary on MNIST

Conv-VAE (2-layer strided convolution encoder, 2-layer transposed convolution decoder) trained on MNIST (100 epochs) shows $AU \approx m$ up to $m = 64$ —no KL-induced saturation. Extending to 300 epochs with fixed β or cyclic KL annealing [3] does not change this ($AU = m$ for $m \leq 64$; FC-VAE saturates at $AU = 15.7$ at $m = 32$ under the same conditions). Extending the m -grid to $\{96, 128, \dots, 256\}$ reveals $AU < m$ onset at $m \approx 96$, but no saturation plateau up to $m = 512$ (Appendix D).

High- β verification. Increasing $\beta_{\max}$ to 10 and 20 shifts the AU saturation to lower dimensions ($m^{\text{knee}} \approx 128$ and ≈ 96 , respectively). This is consistent with the water-filling interpretation: higher β raises $\beta\tau_{\text{eff}}^2$, activating fewer dimensions. The $\beta_{\max} = 4$ failure is therefore not a breakdown of the framework but a sign that $d^*(\beta) \gg 256$ under this setting, requiring stronger regularization or decoder-capacity reduction.

5.8. Natural Images: Defining the Applicability Boundary

Table 7 summarises the key results.

Table 7. Applicability-boundary experiments on CIFAR-10 and SVHN. TwoNN is partially consistent near $m \approx 2\hat{d}_{\text{ID}}$; AU/m^{knee} is not stably identifiable under the tested conditions.

Dataset	Raw TwoNN	Latent TwoNN at $m = 64$	Latent TwoNN at $m = 256$	AU/m^{knee}
CIFAR-10	25.20	25.21 ± 0.41 (consistent)	30.91 (rising)	$AU = m$; boundary
SVHN	10.09	19.97 ± 0.12 (distorted)	27.08 (diverging)	$AU < m$ at $m \geq 128$; unreliable

Extended sweeps ($\beta_{\max}$ up to 80, thin decoder) yield plateau-like AU at $m \approx 512$ on CIFAR-10 (preliminary), but all stable-identification criteria remain unmet. Details in Appendix D.

TwoNN is consistent with the raw-data value near $m \approx 2\hat{d}_{\text{ID}}$ for CIFAR-10, but increases further at larger m (28.83 at $m = 128$, 30.91 at $m = 256$), so the consistency is not uniform across m . SVHN latent TwoNN diverges far from raw TwoNN (10.09 vs. 27.08 at $m = 256$), indicating strong geometric distortion in the latent space.

These experiments **quantify the applicability boundary**: AU/m^{knee} identification on natural images requires substantially stronger regularization and decoder-capacity control than the tested settings provide. Figure 3 provides a visual summary of the diagnostic decisions across all tested settings.

6. Discussion

6.1. Role Separation is the Main Contribution

The core message of this paper is not that m^{knee} estimates effective dimensionality, but that *different indicators serve different diagnostic roles* and conflating them leads to misinterpretation. AU values are reliable monitors of KL-induced activity; TwoNN ID is a stable auxiliary reference (not an unbiased estimator); m^{knee} is informative only under controlled conditions.

6.2. When is m^{knee} Informative?

m^{knee} is informative as a diagnostic aid *only when*: (1) the m -grid is dense near $\hat{d}_{\text{ID}}$; (2) each (m, seed) is initialised independently; (3) at least 3 seeds are used; and (4) AU curves show reasonably monotone behavior (observed in FC-VAE on MNIST). Under these conditions, $m^{\text{knee}} = 10.3 \pm 0.9$ aligns with TwoNN ID (≈ 10),

Figure 4: Applicability-boundary summary
(Green = identifiable, Yellow = conditional, Red = inconclusive/boundary)

Dataset / Model	AU Saturation	TwoNN Consistency	Seed Stability	Diagnostic Decision
MNIST FC-VAE	Yes	Yes ($\sigma \leq 0.12$)	Yes ($\sigma \leq 0.5$)	Identifiable
MNIST Conv-VAE ($\beta=4$)	No (β too low)	Partial ($m=64$ only)	Yes	Inconclusive
MNIST Conv-VAE ($\beta \geq 10$)	Yes ($m_{knee} \approx 128$)	Partial	Partial	Identifiable (high β)
CIFAR-10 (Conv-VAE)	Weak/conditional ($\beta=80$, thin dec.)	Partial ($m \approx 2d$)	No	Boundary
SVHN (Conv-VAE)	Partial	No (diverging)	No	Unreliable

Figure 3. Applicability-boundary summary across all tested settings. Green rows indicate identifiable conditions; yellow rows indicate conditional or inconclusive results; red rows indicate boundary or unreliable regimes. The framework does not only produce a dimension estimate—it also identifies cases where the diagnosis should be treated as inconclusive, which is itself useful knowledge for model selection.

providing a complementary check from the KL-dynamics side. Sparse grids ($m^{\text{knee}} = 43 \pm 15$) are unreliable and should be avoided.

6.3. Optimization Dynamics Interpretation of m^{knee} Instability

Gradient competition hypothesis. The β -ELBO combines reconstruction gradient ($\nabla \mathcal{L}_{\text{rec}}$) and KL gradient ($\beta \nabla \mathcal{L}_{\text{KL}}$). Near $m \approx d^*(\beta)$, these gradients compete: the reconstruction term tries to activate surplus dimensions while the KL term tries to suppress them. The resolution speed depends on the initialization, m -grid, and β , leading to a smooth (not step-like) AU curve in non-linear VAEs and thus grid-dependent m^{knee} values. This hypothesis is consistent with all observed results but has not been directly verified by gradient measurements; it should be regarded as an *interpretive hypothesis*.

6.4. Conv-VAE and Natural Images: Boundary, Not Failure

For Conv-VAEs and natural images, the inability to achieve stable m^{knee} identification is not a failure of the framework—it reveals the *applicability boundary*. The key insight from high- β experiments is that AU saturation does occur when regularization is sufficiently strong relative to decoder capacity. The boundary is thus characterized by the effective τ^2 relative to the data eigenvalue spectrum, consistent with the water-filling interpretation of Analytical Result 4.1.

6.5. Limitations

- **Isometry bias of TwoNN ID.** The reference AE used in Stage 1 has $\kappa \approx 10^9$, so $\hat{d}_{\text{ID}}$ is an approximation. Numerical stability under $m \gg \hat{d}_{\text{ID}}$ is demonstrated, but geometric unbiasedness is not claimed.
- **Single-seed Fashion-MNIST.** Fashion-MNIST is used as a preliminary check only; multi-seed results are not available.
- **Linear-Gaussian approximation.** Analytical Result 4.1 is an analytical model for the linear-Gaussian

surrogate; its applicability to non-linear VAEs is qualitative.

- **Computational cost.** Full m -grid sweeps (11 sizes $\times$ 300 epochs $\times$ 2 models) require hours to days on GPU. IsometricAE Jacobian computation scales as $O(m)$ and is impractical for large m .
- **Natural-image stable identification not achieved.** Under the tested conditions, stable identification on CIFAR-10/SVHN is not achieved; this defines the current applicability boundary.

7. Conclusions

We proposed a multi-indicator diagnostic framework for effective latent dimensionality in AE/VAE representations. The main contribution is not a new dimension estimator but a role separation: AU values as reliable activity monitors, TwoNN/MLE ID as numerically stable auxiliary references, and m^{knee} as an exploratory, grid-dependent change point.

Key findings: (1) AU values and TwoNN ID are reproducible across seeds (MNIST, 3-seed); (2) sparse-grid m^{knee} is unreliable ($\sigma = 15$), whereas dense-grid, independent-initialization, multi-seed evaluation stabilizes it at 10.3 ± 0.9 (TwoNN-consistent); (3) the dominant factor for TwoNN stability is bottleneck margin ($m \gg \hat{d}_{\text{ID}}$), not isometric regularization; (4) Conv-VAE and natural-image experiments quantify the applicability boundary—AU/ m^{knee} diagnostics become unreliable when decoder capacity exceeds the regularization level.

The proposed diagnostic protocol supports practitioners in avoiding over-interpretation of any single indicator, in designing m -grids that make m^{knee} informative, and in recognising when the framework reaches its limits.

Author Contributions: Conceptualization, C.O., M.D., and K.J.; Methodology, K.J.; Software, C.O. and M.D.; Validation, C.O., M.D., and K.J.; Formal Analysis, K.J.; Investigation, C.O. and M.D.; Writing—Original Draft Preparation, K.J.; Writing—Review & Editing, C.O., M.D., and K.J.; Supervision, K.J.; Project Administration, K.J. All authors have read and agreed to the published version of the manuscript.

Institutional Review Board Statement: Not applicable (no human subjects or animals involved).

Informed Consent Statement: Not applicable.

Funding: This research received no external funding.

Data Availability Statement: MNIST [1], Fashion-MNIST [2], CIFAR-10 [4], and SVHN [5] are publicly available benchmark datasets. The source code, configuration files, random seeds, trained-model logs, and numerical results supporting all tables and supplementary analyses are provided as supplementary material for peer review. Upon acceptance, these materials will be deposited in a public repository (GitHub or Zenodo), and the repository URL and DOI will be added to the final version of the manuscript. The repository will include the scripts required to reproduce the AU, TwoNN/MLE, reconstruction-error, and grid-sensitivity analyses reported in this paper.

Conflicts of Interest: The authors declare no conflict of interest.

References

- [1] LeCun, Y.; Bottou, L.; Bengio, Y.; Haffner, P. Gradient-based learning applied to document recognition. *Proc. IEEE* **1998**, *86*, 2278–2324.
- [2] Xiao, H.; Rasul, K.; Vollgraf, R. Fashion-MNIST: A Novel Image Dataset for Benchmarking Machine

Learning Algorithms. *arXiv* **2017**, 1708.07747.

- [3] Fu, H.; Tan, C.; Peng, B.; Zhao, Z.; Yan, T.; Chang, C.; Zhao, L. Cyclical annealing schedule: A simple approach to mitigating KL vanishing. In *Proceedings of NAACL-HLT 2019*, Minneapolis, MN, USA, 2019; pp. 240–250.
- [4] Krizhevsky, A. Learning multiple layers of features from tiny images. *Technical Report*, University of Toronto, 2009.
- [5] Netzer, Y.; Wang, T.; Coates, A.; Bissacco, A.; Wu, B.; Ng, A.Y. Reading digits in natural images with unsupervised feature learning. In *NIPS Workshop on Deep Learning and Unsupervised Feature Learning*, 2011.
- [6] Fefferman, C.; Mitter, S.; Narayanan, H. Testing the manifold hypothesis. *J. Am. Math. Soc.* **2016**, *29*, 983–1049.
- [7] Bengio, Y.; Courville, A.; Vincent, P. Representation learning: A review and new perspectives. *IEEE Trans. Pattern Anal. Mach. Intell.* **2013**, *35*, 1798–1828.
- [8] Facco, E.; d’Errico, M.; Rodriguez, A.; Laio, A. Estimating the intrinsic dimension of datasets by a minimal neighborhood information. *Sci. Rep.* **2017**, *7*, 12140.
- [9] Levina, E.; Bickel, P.J. Maximum likelihood estimation of intrinsic dimension. In *Advances in Neural Information Processing Systems (NIPS)*, Vol. 17, 2005.
- [10] Hatcher, A. *Algebraic Topology*; Cambridge University Press: Cambridge, UK, 2002.
- [11] Kingma, D.P.; Welling, M. Auto-encoding variational Bayes. In *ICLR 2014*, 2014.
- [12] Kato, K.; Zhou, J.; Sasaki, T.; Nakagawa, A. Rate-distortion optimization guided autoencoder for isometric embedding in Euclidean latent space. In *ICML 2020*, Vol. 119, 2020.
- [13] Nazari, P.; Damrich, S.; Hamprecht, F.A. Geometric autoencoders—what you see is what you decode. In *ICML 2023*, Vol. 202, 2023.
- [14] Ceruti, C.; Bassis, S.; Rozza, A.; Lombardi, G.; Casiraghi, E.; Campadelli, P. DANCo: An intrinsic dimensionality estimator exploiting angle and norm concentration. *Pattern Recognit.* **2014**, *47*, 2569–2581.
- [15] Pope, P.; Zhu, C.; Abdelkader, A.; Goldblum, M.; Goldstein, T. The intrinsic dimension of images and its impact on learning. In *ICLR 2021*, 2021.
- [16] Ansuini, A.; Laio, A.; Macke, J.H.; Zoccolan, D. Intrinsic dimension of data representations in deep neural networks. In *NeurIPS 2019*, Vol. 32, 2019.
- [17] Valeriani, L.; Wood, D.; Catania, G.; Gherardini, S.; Laio, A.; Ansuini, A. The geometry of hidden representations of large transformer models. In *NeurIPS 2023*, Vol. 36, 2023.
- [18] Higgins, I.; Matthey, L.; Pal, A.; Burgess, C.; Glorot, X.; Botvinick, M.; Mohamed, S.; Lerchner, A. β -VAE: Learning basic visual concepts with a constrained variational framework. In *ICLR 2017*, 2017.
- [19] Lucas, J.; Tucker, G.; Grosse, R.B.; Norouzi, M. Don’t blame the ELBO! A linear VAE perspective on posterior collapse. In *NeurIPS 2019*, Vol. 32, 2019.

- [20] Rifai, S.; Vincent, P.; Muller, X.; Glorot, X.; Bengio, Y. Contractive auto-encoders: Explicit invariance during feature extraction. In *ICML 2011*, Vol. 28, 2011.
- [21] Alain, G.; Bengio, Y. What regularized auto-encoders learn from the data-generating distribution. *J. Mach. Learn. Res.* **2014**, *15*, 3743–3773.

A. Single-Seed VAE Illustration

Table 8 shows the single-seed preliminary run that motivated the multi-seed study. The $m^{\text{knee}} = 12$ value observed here depends on the specific grid ($m \in \{1, 4, 8, \dots, 256\}$) and initialization, and is *not* used as a main result.

Table 8. MNIST VAE ($\beta = 4$, 300 epochs, seed 42, extended grid) — illustration only. The slowdown at $m = 12$ and the non-monotone AU (29 at $m = 64$, 26 at $m = 128$) reflect stochastic optimization, not a stable signal. The main stabilized result is $m^{\text{knee}} = 10.3 \pm 0.9$ (Table 5).

m	MSE	AU	TwoNN ID	Note
1	0.0464	1	1.00	
8	0.0186	8	6.36	increasing
12	0.0184	8	6.26	ρ drops (single-seed, not main result)
64	0.0103	29	9.16	
128	0.0085	26	10.02	non-monotone AU
256	0.0072	36	10.14	

B. Proofs and Supplementary Theory

B.1. Topological Lower Bound

Proposition B.1 (Topological Lower Bound). *To represent a smooth d -dimensional manifold $\mathcal{M} \subset \mathbb{R}^N$ without self-intersections under zero reconstruction error (idealised setting), a bottleneck of $m \geq d_{\text{emb}}$ is necessary, where d_{emb} is the topological minimum embedding dimension. For simply connected manifolds (e.g. Swiss Roll), $d_{\text{emb}} = d_{\text{ID}}$. For the Torus T^2 , $d_{\text{ID}} = 2$ but $d_{\text{emb}} = 3$: it cannot be embedded in $\mathbb{R}^2$ without self-intersections, explaining the MSE elbow at $m = 3$ (see [10], Ch. 2).*

Note: This is an idealised result; finite capacity and approximate reconstruction are excluded.

B.2. Proof Sketch for Analytical Result 4.1

The derivation follows the standard linear-Gaussian rate-distortion analysis with a β -ELBO objective. For each PCA direction j , the contribution $L_j = -D_j/(2\tau^2) - \beta R_j$ (rate-distortion surrogate) is maximised independently. The activation threshold follows a two-case analysis:

Case 1 ($\lambda_j \leq \beta\tau^2$, **inactive**). The water-filling solution $D_j^* = \beta\tau^2$ is infeasible (requires $D_j^* \leq \lambda_j$), so the optimum is the inactive solution with zero rate ($R_j = 0$, $D_j = \lambda_j$).

Case 2 ($\lambda_j > \beta\tau^2$, **potentially active**). The water-filling solution is feasible. The gain from activating direction j is $\Delta L_j = L_j^{\text{active}} - L_j^{\text{dead}} = (\beta/2) h(\lambda_j/\beta\tau^2)$ where $h(t) = t - 1 - \log t$. Since $h(t) > 0$ for all $t > 0$, $t \neq 1$, and here $t = \lambda_j/\beta\tau^2 > 1$, we have $\Delta L_j > 0$: activating is strictly better than inactivating.

Combining both cases: direction j is active if and only if $\lambda_j > \beta\tau^2$.

Remark B.2 (KL vs. mutual information). *The surrogate L_j implicitly identifies the expected KL $\mathbb{E}_x \text{KL}(q_j \| p_j)$*

with the mutual information $R_j = I(z_j; y_j)$. For a true VAE, the expected KL also includes aggregated-posterior/prior mismatch. Therefore Analytical Result 4.1 is a result for the linear-Gaussian rate-distortion surrogate, not a rigorous theorem about the full VAE ELBO.

C. TwoNN Stability Details and Linear Theory Bridge

C.1. Exp-Geom: MNIST AE m -sweep

Table 9. MNIST AE m -sweep (Exp-Geom). TwoNN stabilizes at $m \geq 48 \approx 5\hat{d}_{\text{ID}}$.

m	κ	TwoNN ID	Trustworthiness	Convergence
4	2.69	4.21	0.960	not converged
16	3.28	8.17	0.993	converging
48	4.44	10.16	0.998	converged
64	4.25	10.17	0.998	stable
96	9.94	10.86	0.999	stable (despite κ jump)
128	37.0	10.98	0.999	stable
$n_{\text{train}} = 10\,000$, 200 ep, seed=42.				

C.2. Exp-Iso-ID: Standard AE vs. IsometricAE on Swiss Roll

Table 10. Latent TwoNN ID for AE and IsometricAE on Swiss Roll ($d_{\text{true}} = 2$, input TwoNN= 2.427). At $m \geq 3$, TwoNN difference stays below $\pm 2.3\%$ despite κ ratio ~ 2 .

m	AE		IsometricAE		Diff. (%)
	TwoNN ID	κ	TwoNN ID	κ	
2	1.780	1.76	1.799	1.08	+1.1
3	2.458	2.10	2.454	1.11	−0.2
4	2.483	1.80	2.479	1.14	−0.2
6	2.421	1.79	2.474	1.10	+2.2
8	2.500	1.45	2.443	1.11	−2.3

C.3. Fashion-MNIST Exp-Real-Geom (Preliminary)

Fashion-MNIST ($\hat{d}_{\text{ID}} = 9$, $m \in \{9, 18, 27, 45, 90\}$) shows qualitatively the same pattern as MNIST: TwoNN converges at $m \approx 5\hat{d}_{\text{ID}}$, inter-model differences are small, and κ improvements have minimal effect. The Fashion-MNIST experiment is used as a preliminary single-seed external check; multi-seed validation is left for future work and is not used to support the main claims.

C.4. Linear Theory Bridge (Exp-Mnist-PCA)

MNIST PCA spectrum: $\lambda_1 = 5.22, \lambda_2 = 3.76, \dots, \lambda_{10} = 1.23$. At $\beta = 4$ and back-calculated $\tau_{\text{eff}}^2 \approx 0.30$, the water-filling prediction gives $d^* \approx 10$, consistent with TwoNN ID. The smooth (not step-like) AU curve of the non-linear FC-VAE explains why sparse-grid m^{knee} is unstable: gradual saturation violates the step-function assumption of Proposition 4.4.

D. Conv-VAE and Natural-Image Details

D.1. Exp-Conv-M512: MNIST Conv-VAE Extended Grid

Table 11. MNIST Conv-VAE with $m \in \{256, 384, 512\}$ (per seed; KL cyclic annealing, $\beta_{\max} = 4, 150$ ep). No AU plateau up to $m = 512$; $d^*(\beta) > 512$ under this setting.

m	AU (s42)	AU (s123)	AU (s777)
256	241	245	238
384	382	383	380
512	512	511	512

D.2. Natural-Image Extended Sweep (Exp-NatImg-Sweep)

Extended sweeps on CIFAR-10 ($m \in [32, 1024]$, $\beta_{\max} \in \{10, 20, 40, 80\}$, Standard/Deep/Thin decoders, 300 ep, 3 seeds):

- $\beta_{\max} = 80$, Thin decoder: AU plateau-like at $m \in [512, 1024]$ ($AU \approx 24$). This is a preliminary, conditional observation.
- TwoNN ID: 24.7–25.4 ($\sigma \leq 0.6$) across all conditions.

The five stable-identification criteria (seed stability, grid stability, β stability, AU-plateau stability, TwoNN–AU consistency) are all unmet under the tested conditions; strict stable identification remains a future goal.

E. Sensitivity Analyses

E.1. AU Threshold δ

Swiss Roll and Torus experiments confirm AU patterns are identical for $\delta \in \{0.001, 0.01, 0.05, 0.1\}$ (active coordinate variance $O(1)$ vs. inactive $O(10^{-4})$; 2-order gap).

E.2. τ_{knee} Sensitivity

Cross-dataset validation confirms $\tau = 0.25$ is stable for Gaussian mixture data. Fashion-MNIST shows boundary sensitivity near $\tau \approx 0.25$; the second-difference alternative is recommended for gradual-slope curves. Dense-grid experiments confirm the fixed- τ method outperforms second differences ($\sigma = 0.9$ vs. 18.3 ± 6.8) for MNIST.

E.3. β Sensitivity

Preliminary 50-epoch MNIST subsets: $\beta = 1$ gives $AU \approx m$ (insufficient regularization); $\beta = 8$ shows over-suppression signs. AU patterns are stable for $\beta \in [2, 8]$; $\beta = 4$ is used as default.

Publisher’s Note: MDPI stays neutral with regard to jurisdictional claims in published maps and institutional affiliations.

© 2025 by the authors. Licensee MDPI, Basel, Switzerland. This article is an open access article distributed under the terms and conditions of the Creative Commons Attribution (CC BY) license (<https://creativecommons.org/licenses/by/4.0/>).